

Avoiding tipping points in the management of ecological systems: a non-parametric Bayesian approach

Carl Boettiger^{a,*}, Marc Mangel^a, Stephan B. Munch^b

^a*Center for Stock Assessment Research, Department of Applied Math and Statistics, University of California, Mail Stop SOE-2, Santa Cruz, CA 95064, USA*

^b*Southwest Fisheries Science Center, National Oceanic and Atmospheric Administration, 110 Shaffer Road, Santa Cruz, CA 95060, USA*

```
options(xtable.print.comment = FALSE)
options(xtable.type = "latex", table.placement = "H")
opts_knit$set(progress = TRUE, verbose = TRUE)
opts_chunk$set(dev = c("pdf", "png"), fig.width = 5.5, fig.height = 4, cache.path = "cache/nonparametric-1",
  cache = TRUE, external = TRUE)
opts_chunk$set(warning = FALSE, message = FALSE, comment = NA, tidy = FALSE)
toggle = "markup"
theme_set(theme_bw(base_size = 12))

require(modeest)
posterior.mode <- function(x) {
  mlv(x, method="shorth")$M
}

f <- RickerAllee
p <- c(2, 8, 5)
K <- 10 # approx, a li'l' less
allee <- 5 # approx, a li'l' less

sigma_g <- 0.05
sigma_m <- 0.0
z_g <- function() rlnorm(1, 0, sigma_g)
z_m <- function() 1
x_grid <- seq(0, 1.5 * K, length=50)
```

*Corresponding author
Email address: cboettig@ucsc.edu (Carl Boettiger)

```

h_grid <- x_grid
profit <- function(x,h) pmin(x, h)
delta <- 0.01
OptTime <- 50 # stationarity with unstable models is tricky thing
reward <- 0
xT <- 0
Xo <- allee+.5# observations start from
x0 <- K # simulation under policy starts from
Tobs <- 40
MaxT = 1000 # timeout for value iteration convergence

set.seed(1234)
#harvest <- sort(rep(seq(0, .5, length=7), 5))
x <- numeric(Tobs)
x[1] <- Xo
nz <- 1
for(t in 1:(Tobs-1))
  x[t+1] = z_g() * f(x[t], h=0, p=p)
obs <- data.frame(x = c(rep(0,nz),
                        pmax(rep(0,Tobs-1), x[1:(Tobs-1)])),
                  y = c(rep(0,nz),
                        x[2:Tobs]))

raw_plot <- ggplot(data.frame(time = 1:Tobs, x=x), aes(time,x)) + geom_line()
raw_plot

set.seed(12345)
estf <- function(p){
  mu <- f(obs$x,0,p)
  -sum(dlnorm(obs$y, log(mu), p[4]), log=TRUE)
}
par <- c(p[1]*rlnorm(1,0,.1),
        p[2]*rlnorm(1,0,.1),
        p[3]*rlnorm(1,0,.1),
        sigma_g * rlnorm(1,0,.1))
o <- optim(par, estf, method="L", lower=c(1e-5,1e-5,1e-5,1e-5))
f_alt <- f

```




Figure 1: plot of chunk obs

```
p_alt <- c(as.numeric(o$par[1]), as.numeric(o$par[2]), as.numeric(o$par[3]))
sigma_g_alt <- as.numeric(o$par[4])
```

```
est <- list(f = f_alt, p = p_alt, sigma_g = sigma_g_alt, mloglik=o$value)
```

```
true_means <- sapply(x_grid, f, 0, p)
est_means <- sapply(x_grid, est$f, 0, est$p)
```

```
#inv gamma has mean b / (a - 1) (assuming a>1) and variance b ^ 2 / ((a - 2) * (a - 1) ^ 2) (assuming a>2)
s2.p <- c(5,5)
d.p = c(10, 1/0.1)
```

```
gp <- gp_mcmc(obs$x, y=obs$y, n=1e5, s2.p = s2.p, d.p = d.p)
gp_dat <- gp_predict(gp, x_grid, burnin=1e4, thin=300)
```

```
gp_assessment_plots <- summary_gp_mcmc(gp, burnin=1e4, thin=300)
```

```
# Summarize the GP model
tgp_dat <-
```

```

data.frame( x = x_grid,
            y = gp_dat$E_Ef,
            ymin = gp_dat$E_Ef - 2 * sqrt(gp_dat$E_Vf),
            ymax = gp_dat$E_Ef + 2 * sqrt(gp_dat$E_Vf) )

y <- x
N <- length(x);
jags.data <- list("N","y")
n.chains <- 6
n.iter <- 1e6
n.burnin <- floor(10000)
n.thin <- max(1, floor(n.chains * (n.iter - n.burnin)/1000))
n.update <- 10

stdQ_prior_p <- c(1e-6, 100)
stdR_prior_p <- c(1e-6, .1)
stdQ_prior <- function(x) dunif(x, stdQ_prior_p[1], stdQ_prior_p[2])
stdR_prior <- function(x) dunif(x, stdR_prior_p[1], stdR_prior_p[2])

K_prior_p <- c(0.01, 20.0)
r0_prior_p <- c(0.01, 6.0)
theta_prior_p <- c(0.01, 20.0)

bugs.model <-
paste(sprintf(
"model{
  K      ~ dunif(%s, %s)
  r0     ~ dunif(%s, %s)
  theta ~ dunif(%s, %s)
  stdQ ~ dunif(%s, %s)",
  K_prior_p[1], K_prior_p[2],
  r0_prior_p[1], r0_prior_p[2],
  theta_prior_p[1], theta_prior_p[2],
  stdQ_prior_p[1], stdQ_prior_p[2]),
"
iQ <- 1 / (stdQ * stdQ);

```

```

y[1] ~ dunif(0, 10)
for(t in 1:(N-1)){
  mu[t] <- log(y[t]) + r0 * (1 - y[t]/K)* (y[t] - theta) / K
  y[t+1] ~ dlnorm(mu[t], iQ)
}
})"
writeLines(bugs.model, "allen_process.bugs")

K_prior      <- function(x) dunif(x, K_prior_p[1], K_prior_p[2])
r0_prior <- function(x) dunif(x, r0_prior_p[1], r0_prior_p[2])
theta_prior <- function(x) dunif(x, theta_prior_p[1], theta_prior_p[2])
par_priors  <- list(K = K_prior, deviance = function(x) 0 * x,
                   r0 = r0_prior, theta = theta_prior,
                   stdQ = stdQ_prior)

jags.params=c("K","r0","theta","stdQ") # be sensible about the order here
jags.inits <- function(){
  list("K"= 10 * rlnorm(1,0, 0.1),
       "r0"= 1 * rlnorm(1,0, 0.1) ,
       "theta"= 5 * rlnorm(1,0, 0.1) ,
       "stdQ"= abs( 0.1 * rlnorm(1,0, 0.1)),
       .RNG.name="base::Wichmann-Hill", .RNG.seed=123)
}

set.seed(1234)
# parallel refuses to take variables as arguments (e.g. n.iter = 1e5 works, but n.iter = n doesn't)
allen_jags <- do.call(jags.parallel, list(data=jags.data, inits=jags.inits,
                                         jags.params, n.chains=n.chains,
                                         n.iter=n.iter, n.thin=n.thin,
                                         n.burnin=n.burnin,
                                         model.file="allen_process.bugs"))

# Run again iteratively if we haven't met the Gelman-Rubin convergence criterion
recompile(allen_jags) # required for parallel

```

Compiling model graph
 Resolving undeclared variables

Allocating nodes

Graph Size: 328

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 328

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 328

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 328

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 328

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 328

Initializing model

```
allen_jags <- do.call(autojags,
                      list(object=allen_jags, n.update=n.update,
                           n.iter=n.iter, n.thin = n.thin))

tmp <- lapply(as.mcmc(allen_jags), as.matrix) # strip classes the hard way...
allen_posteriors <- melt(tmp, id = colnames(tmp)[[1]])
names(allen_posteriors) = c("index", "variable", "value", "chain")
plot_allen_traces <- ggplot(allen_posteriors) + geom_line(aes(index, value)) +
  facet_wrap(~ variable, scale="free", ncol=1)

allen_priors <- ddply(allen_posteriors, "variable", function(dd){
  grid <- seq(min(dd$value), max(dd$value), length = 100)
  data.frame(value = grid, density = par_priors[[dd$variable[1]]](grid))
})
plot_allen_posteriors <- ggplot(allen_posteriors, aes(value)) +
  stat_density(geom="path", position="identity", alpha=0.7) +
  geom_line(data=allen_priors, aes(x=value, y=density), col="red") +
  facet_wrap(~ variable, scale="free", ncol=3)

A <- allen_posteriors
A$index <- A$index + A$chain * max(A$index) # Combine samples across chains by renumbering index
pardist <- acast(A, index ~ variable)
bayes_coef <- apply(pardist, 2, posterior.mode)
bayes_pars <- unname(c(bayes_coef["r0"], bayes_coef["K"], bayes_coef["theta"])) # parameters formatted for
allen_f <- function(x,h,p) unname(RickerAllee(x,h, unname(p[c("r0", "K", "theta")])))
allen_means <- sapply(x_grid, f, 0, bayes_pars)
bayes_pars

[1] 0.5003 7.7115 4.6039

head(pardist)

      K deviance      r0      stdQ theta
```

```

170 7.898    34.41 0.6668 0.05566 3.956
171 7.745    32.76 0.7554 0.05137 4.668
172 7.577    31.66 1.0705 0.04562 4.073
173 7.859    31.26 1.5647 0.04238 5.087
174 7.821    30.67 1.7348 0.04662 5.230
175 8.003    34.90 1.4495 0.05005 5.444

```

```

K_prior_p <- c(0.01, 40.0)
r0_prior_p <- c(0.01, 20.0)
bugs.model <-
paste(sprintf(
"model{
  K    ~ dunif(%s, %s)
  r0    ~ dunif(%s, %s)
  stdQ ~ dunif(%s, %s)",
  K_prior_p[1], K_prior_p[2],
  r0_prior_p[1], r0_prior_p[2],
  stdQ_prior_p[1], stdQ_prior_p[2]),
"
  iQ <- 1 / (stdQ * stdQ);
  y[1] ~ dunif(0, 10)
  for(t in 1:(N-1)){
    mu[t] <- log(y[t]) + r0 * (1 - y[t]/K)
    y[t+1] ~ dlnorm(mu[t], iQ)
  }
}")
writeLines(bugs.model, "ricker_process.bugs")

K_prior    <- function(x) dunif(x, K_prior_p[1], K_prior_p[2])
r0_prior   <- function(x) dunif(x, r0_prior_p[1], r0_prior_p[2])
par_priors <- list(K = K_prior, deviance = function(x) 0 * x,
                  r0 = r0_prior, stdQ = stdQ_prior)

jags.params=c("K","r0", "stdQ")
jags.inits <- function(){
  list("K"= 10 * rlnorm(1,0,.5),
       "r0"= rlnorm(1,0,.5),

```



```

        "stdQ"=sqrt(0.05) * rlnorm(1,0,.5),
        .RNG.name="base::Wichmann-Hill", .RNG.seed=123)
}
set.seed(12345)
ricker_jags <- do.call(jags.parallel,
                      list(data=jags.data, inits=jags.inits,
                           jags.params, n.chains=n.chains,
                           n.iter=n.iter, n.thin=n.thin, n.burnin=n.burnin,
                           model.file="ricker_process.bugs"))
recompile(ricker_jags)

```

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 249

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 249

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 249

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 249

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 249

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 249

Initializing model

```
ricker_jags <- do.call(autojags,  
                      list(object=ricker_jags, n.update=n.update,  
                          n.iter=n.iter, n.thin = n.thin,  
                          progress.bar="none"))
```

```
tmp <- lapply(as.mcmc(ricker_jags), as.matrix) # strip classes the hard way...  
ricker_posteriors <- melt(tmp, id = colnames(tmp)[[1]])  
names(ricker_posteriors) = c("index", "variable", "value", "chain")  
plot_ricker_traces <- ggplot(ricker_posteriors) + geom_line(aes(index, value)) +  
  facet_wrap(~ variable, scale="free", ncol=1)
```

```
ricker_priors <- ddply(ricker_posteriors, "variable", function(dd){  
  grid <- seq(min(dd$value), max(dd$value), length = 100)  
  data.frame(value = grid, density = par_priors[[dd$variable[1]]](grid))  
})
```

plot posterior distributions

```
plot_ricker_posteriors <- ggplot(ricker_posteriors, aes(value)) +  
  stat_density(geom="path", position="identity", alpha=0.7) +  
  geom_line(data=ricker_priors, aes(x=value, y=density), col="red") +  
  facet_wrap(~ variable, scale="free", ncol=2)
```

```

A <- ricker_posteriors
A$index <- A$index + A$chain * max(A$index) # Combine samples across chains by renumbering index
ricker_pardist <- acast(A, index ~ variable)
bayes_coef <- apply(ricker_pardist, 2, posterior.mode)
ricker_bayes_pars <- unname(c(bayes_coef["r0"], bayes_coef["K"]))
ricker_f <- function(x, h, p){
  sapply(x, function(x){
    x <- pmax(0, x-h)
    pmax(0, x * exp(p["r0"] * (1 - x / p["K"] )) )
  })
}
ricker_means <- sapply(x_grid, Ricker, 0, ricker_bayes_pars[c(1,2)])
head(ricker_pardist)

```

		K	deviance	r0	stdQ
170	7.727	35.73	0.29440	0.05197	
171	7.949	37.93	0.18845	0.05872	
172	32.183	44.54	0.02474	0.05688	
173	8.149	37.59	0.12204	0.05407	
174	7.460	39.01	0.38643	0.05495	
175	7.361	35.54	0.23805	0.05150	

```
ricker_bayes_pars
```

```
[1] 0.1997 7.5976
```

```

r0_prior_p <- c(.0001, 10.0)
theta_prior_p <- c(.0001, 10.0)
K_prior_p <- c(.0001, 40.0)
bugs.model <-
paste(sprintf(
"model{
  r0    ~ dunif(%s, %s)
  theta ~ dunif(%s, %s)
  K     ~ dunif(%s, %s)
  stdQ ~ dunif(%s, %s)",
  r0_prior_p[1], r0_prior_p[2],

```

```

theta_prior_p[1], theta_prior_p[2],
K_prior_p[1], K_prior_p[2],
stdQ_prior_p[1], stdQ_prior_p[2]),

"

iQ <- 1 / (stdQ * stdQ);

y[1] ~ dunif(0, 10)
for(t in 1:(N-1)){
  mu[t] <- log(r0) + theta * log(y[t]) - log(1 + pow(abs(y[t]), theta) / K)
  y[t+1] ~ dlnorm(mu[t], iQ)
}
}"))
writeLines(bugs.model, "myers_process.bugs")

K_prior      <- function(x) dunif(x, K_prior_p[1], K_prior_p[2])
r_prior      <- function(x) dunif(x, r0_prior_p[1], r0_prior_p[2])
theta_prior  <- function(x) dunif(x, theta_prior_p[1], theta_prior_p[2])
par_priors <- list( deviance = function(x) 0 * x, K = K_prior,
                   r0 = r_prior, theta = theta_prior,
                   stdQ = stdQ_prior)

jags.params=c("r0", "theta", "K", "stdQ")
jags.inits <- function(){
  list("r0"= 1 * rlnorm(1,0,.1),
       "K"= 10 * rlnorm(1,0,.1),
       "theta" = 1 * rlnorm(1,0,.1),
       "stdQ"= sqrt(0.2) * rlnorm(1,0,.1),
       .RNG.name="base::Wichmann-Hill", .RNG.seed=123)
}

set.seed(12345)
myers_jags <- do.call(jags.parallel,
                     list(data=jags.data, inits=jags.inits,
                          jags.params, n.chains=n.chains,
                          n.iter=n.iter, n.thin=n.thin,
                          n.burnin=n.burnin,
                          model.file="myers_process.bugs"))

```

```
recompile(myers_jags)
```

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 406

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 406

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 406

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 406

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 406

Initializing model

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 406

Initializing model

```
myers_jags <- do.call(autojags,
                        list(myers_jags, n.update=n.update,
                             n.iter=n.iter, n.thin = n.thin,
                             progress.bar="none"))

tmp <- lapply(as.mcmc(myers_jags), as.matrix) # strip classes
myers_posteriors <- melt(tmp, id = colnames(tmp[[1]]))
names(myers_posteriors) = c("index", "variable", "value", "chain")
plot_myers_traces <- ggplot(myers_posteriors) + geom_line(aes(index, value)) +
  facet_wrap(~ variable, scale="free", ncol=1)

par_prior_curves <- ddply(myers_posteriors, "variable", function(dd){
  grid <- seq(min(dd$value), max(dd$value), length = 100)
  data.frame(value = grid, density = par_priors[[dd$variable[1]]](grid))
})

plot_myers_posteriors <- ggplot(myers_posteriors, aes(value)) +
  stat_density(geom="path", position="identity", alpha=0.7) +
  geom_line(data=par_prior_curves, aes(x=value, y=density), col="red") +
  facet_wrap(~ variable, scale="free", ncol=3)

A <- myers_posteriors
A$index <- A$index + A$chain * max(A$index) # Combine samples across chains by renumbering index
myers_pardist <- acast(A, index ~ variable)
bayes_coef <- apply(myers_pardist, 2, posterior.mode) # much better estimates
myers_bayes_pars <- unname(c(bayes_coef[2], bayes_coef[3], bayes_coef[1]))
myers_means <- sapply(x_grid, Myer_harvest, 0, myers_bayes_pars)
myers_f <- function(x,h,p) Myer_harvest(x, h, p[c("r0", "theta", "K")])
head(myers_pardist)

      K deviance      r0      stdQ  theta
```

```

170 20.51    37.16 1.3995 0.04136 0.9864
171 17.10    41.09 0.6076 0.05170 1.8684
172 24.40    36.33 1.1692 0.04307 1.0715
173 27.56    37.46 0.9533 0.05550 1.2000
174 22.21    34.94 0.5140 0.04203 1.8694
175 36.05    33.49 0.3089 0.03977 2.1519

```

```
myers_bayes_pars
```

```
[1] 35.4576 0.5129 23.9914
```

```

models <- data.frame(x=x_grid,
                     GP=tdp_dat$y,
                     True=true_means,
                     MLE=est_means,
                     Ricker=ricker_means,
                     Allen = allen_means,
                     Myers = myers_means)
models <- melt(models, id="x")

# some labels
names(models) <- c("x", "method", "value")

# labels for the colorkey too
model_names = c("GP", "True", "MLE", "Ricker", "Allen", "Myers")
colorkey=cbPalette
names(colorkey) = model_names

# uses expected values from GP, instead of integrating over posterior
#matrices_gp <- gp_transition_matrix(gp_dat$E_Ef, gp_dat$E_Vf, x_grid, h_grid)
matrices_gp <- gp_transition_matrix(gp_dat$Ef_posterior, gp_dat$Vf_posterior, x_grid, h_grid)
opt_gp <- value_iteration(matrices_gp, x_grid, h_grid, MaxT, xT, profit, delta, reward)

matrices_true <- f_transition_matrix(f, p, x_grid, h_grid, sigma_g)
opt_true <- value_iteration(matrices_true, x_grid, h_grid, OptTime=MaxT, xT, profit, delta=delta)
matrices_estimated <- f_transition_matrix(est$f, est$p, x_grid, h_grid, est$sigma_g)
opt_estimated <- value_iteration(matrices_estimated, x_grid, h_grid, OptTime=MaxT, xT, profit, delta=delta)

```

```

matrices_allen <- parameter_uncertainty_SDP(allen_f, x_grid, h_grid, pardist, 4)
opt_allen <- value_iteration(matrices_allen, x_grid, h_grid, OptTime=MaxT, xT, profit, delta=delta)

matrices_ricker <- parameter_uncertainty_SDP(ricker_f, x_grid, h_grid, as.matrix(ricker_pardist), 3)
opt_ricker <- value_iteration(matrices_ricker, x_grid, h_grid, OptTime=MaxT, xT, profit, delta=delta)

matrices_myers <- parameter_uncertainty_SDP(myers_f, x_grid, h_grid, as.matrix(myers_pardist), 4)
myers_alt <- value_iteration(matrices_myers, x_grid, h_grid, OptTime=MaxT, xT, profit, delta=delta)

OPT = data.frame(GP = opt_gp$D, True = opt_true$D, MLE = opt_estimated$D, Ricker = opt_ricker$D, Allen = 
colorkey=cbPalette
names(colorkey) = names(OPT)

sims <- lapply(OPT, function(D){
  set.seed(1)
  lapply(1:100, function(i)
    ForwardSimulate(f, p, x_grid, h_grid, x0, D, z_g, profit=profit, OptTime=OptTime)
  )
})

dat <- melt(sims, id=names(sims)[[1]][[1]])
sims_data <- data.table(dat)
setnames(sims_data, c("L1", "L2"), c("method", "reps"))
# Legend in original ordering please, not alphabetical:
sims_data$method = factor(sims_data$method, ordered=TRUE, levels=names(OPT))

Profit <- sims_data[, sum(profit), by=c("reps", "method")]
tmp <- dcast(Profit, reps ~ method)
#tmp$Allen <- tmp[, "Allen"] + rnorm(dim(tmp)[1], 0, 1) # jitter for plotting
tmp <- tmp / tmp[, "True"]
tmp <- melt(tmp[2:dim(tmp)[2]])
actual_over_optimal <- subset(tmp, variable != "True")

allen_deviance <- posterior.mode(pardist[, 'deviance'])
ricker_deviance <- posterior.mode(ricker_pardist[, 'deviance'])
myers_deviance <- posterior.mode(myers_pardist[, 'deviance'])
true_deviance <- 2*estf(c(p, sigma_g))

```



```
mle_deviance <- 2*estf(c(est$p, est$sigma_g))
```

```
c(allen = allen_deviance, ricker=ricker_deviance, myers=myers_deviance, true=true_deviance, mle=mle_deviance)
```

allen	ricker	myers	true	mle
32.93	35.92	35.46	-61.08	-889.99

Abstract

Decision-theoretic methods often rely on simple parametric models of ecological dynamics to compare the value of a potential sequence of actions. Unfortunately, such simple models rarely capture the complexity or uncertainty found in most real ecosystems. Non-parametric Bayesian methods offer a promising statistical approach for predictive modeling of ecological dynamics in regions of state space where the data is adequate, while at the same time offering more flexible patterns with greater uncertainty outside the observed data. This contrasts from simple parametric models which provide relatively constant level of uncertainty in regions with and without adequate data. The consequence of such misplaced confidence outside the data can lead to highly undesirable results that may be avoided with the more flexible non-parametric Bayesian approach.

Introduction

Most management recommendations from the ecological literature are based on (or at least motivated by) parametric models. Though in principle these models reflect mechanistic underpinnings of biological interactions involved, in practice real mechanisms are often unknown and the true dynamics too complex to be captured by simple models. While simple mechanistic models can nevertheless provide important insights into possible dynamics – for instance, demonstrating that a vaccine does not have to be 100% effective to eliminate the transmission of a virus – such approaches are not well-suited for use in forecasting outcomes of potential management options. Non-parametric approaches offer a more flexible alternative that can both more accurately reflect the data available while also representing greater uncertainty in areas (of state-space) where data is lacking. Ecological research and management strategy should pay closer attention to the opportunities and challenges nonparametric modeling can offer.

Quantitative vs Qualitative Decisions

In this paper, we consider those ecological management problems in which a mathematical (or computational) model is used to quantitatively inform decision-making by allowing a manager to compare to the expected consequences of potential management actions (or policies). We distinguish this from the solely qualitative use of

a model, in which models are used to represent and compare hypotheses of different mechanisms that would lead to qualitatively different actions.

Though simple mechanistically motivated models may be best for the latter case (@Cuddington2013; Geritz and Kisdi 2012), such models can be not only inaccurate but misleading in quantitative decision making. This arises for two reasons:

1. We do not know what the correct parameter values are for the models.
2. We do not know what the correct models are for ecological systems.

The complexity of ecological interactions and a lack of data contribute greatly to both of the problems. This concern is particularly acute in case of ecological tipping points (Barnosky et al. 2012; Scheffer et al. 2001), which arise from feedbacks common in ecological systems and can lead to a sudden catastrophic transition to an undesirable state. In general we do not know where such tipping points are unless we have already observed the transition, in which case it is frequently too late to respond. Though there may exist early warning signals for certain kinds of these transitions that are driven by slow changes (Scheffer et al. 2009), we do not know when, where, or how to apply them to the decision making context more generally (Boettiger and Hastings 2013).

In addition to facing the potentially grave consequences of such transitions, if they do exist in a system of interest, this generally means that we lack data in the region near and beyond a possible tipping point. As a consequence, our models are least accurate just where we need them to be most accurate.

Unfortunately, parametric modeling approaches are not well-suited to this scenario, in which good data is available over only part of the relevant state-space. Necessary assumptions about model structure in any parametric approach can lead such models to perform very poorly, even when accounting for parameter uncertainty.

Non-parametric models can more realistically represent uncertainties outside of observed data while also better capturing the dynamics in the region observed. Advances in the theory and computational implementations of nonparametric methods make them ripe for such applications. We use the classic problem of optimal harvest of a marine fishery to illustrate how the nonparametric approach of Gaussian processes can be applied. We will compare Bayesian implementations of both nonparametric and parametric models, which best allow us to capture the uncertainty of model estimation in either case and permits a more natural comparison between approaches.

Terminology: “Non-parametric”

The term non-parametric to describe modeling approaches such as Gaussian processes is a common source of confusion, notwithstanding the fact that the model is still specified by parameters. Some literature has introduced the term “semi-parametric” in tacit acknowledgment of this, which no doubt only contributes to the confusion. The problem is further exacerbated by the several meanings assigned to the term in statistics: (a) that the method does not assume some particular probability distribution (e.g. a non-parametric bootstrap) or (b) that

the method does not assume a fixed structure to the model. Our use is that of (b), in contrast to the classical approaches that do. The difference is clearest by way of example.

Uncertainties: the classical approach

The ecological dynamics of most systems of management interest are typically both complex and unknown. Despite this, quantitative policy optimization in these systems has almost exclusively been based on simple, mechanistically motivated parametric models. There are several reasons for this. First, limited available data almost always precludes any accurate estimate of more complicated models. Second, the computational demands of determining an optimal strategy in a sequential decision-theory problem suffer from the “curse of dimensionality” that makes it difficult to consider more complicated models, states, or actions. Unfortunately, these challenges also make it particularly difficult to reflect the true uncertainty either in the parameters of the given models (parametric uncertainty) or in the structure of the models themselves (structural uncertainty). Further unknowns such as measurement uncertainty, parameter uncertainty, unobserved states, knowledge of boundary conditions, etc. further compound the issue. Though a hierarchical Bayesian approach provides a natural way to address these from a statistical standpoint, formulating reasonable parametric descriptions of each form of uncertainty is a challenging task in itself, let alone the computational difficulties of solving such a system. Cressie et al. (2009) provides a good account of the successes and challenges of the approach. Applying these approaches in the management context of sequential decision making, in which forecasts must be obtained over a range of possible actions and updated regularly as new information arrives makes such an approach less feasible still.

An active literature and growing computational power over the past several decades have only marginally improved this situation. Parametric or structural uncertainty can be introduced only by increasing the state space to reflect a distribution of possible parameters or the degree of belief in each of a set of possible models. Frequently, the space of possible actions must then be reduced or the algorithms adjusted by approximations to keep computations feasible.

Traditional approaches to optimal control (Pontryagin’s principle, stochastic dynamic programming) rely on knowledge of the state equation, usually described by a simple parametric model. Here we illustrate how a stochastic dynamic programming algorithm can alternatively be driven by the predictions from a Gaussian process.

Approach and Methods

Underlying Model

Concerns over the potential for tipping points in ecological dynamics (Scheffer et al. 2001) highlight the dangers of uncertainty in ecological management and pose a substantial challenge to existing decision-theoretic approaches (Brozović and Schlenker 2011). Because intervention is often too late after a tipping point has been crossed (but

see Hughes et al. (2013)), management is most often concerned with avoiding potentially catastrophic tipping points before any data is available at or following a transition that would more clearly reveal these regime shift dynamics (e.g. Bestelmeyer et al. 2012).

To illustrate the value of the non-parametric Bayesian approach to management, we focus on example of a system containing such a tipping point whose dynamics can still be described by a simple, one-dimensional parametric model.

We will focus on a simple parametric model for a single species (derived from first principles by Allen et al. 2005) as our underlying “reality”.

$$X_{t+1} = Z_t f(S_t) \quad (1)$$

$$S_t = X_t - h_t \quad (2)$$

$$f(S_t) = S_t e^{r(1-\frac{S_t}{K})(S_t-C)} \quad (3)$$

Where Z_t is multiplicative noise function with mean 1, representing stochastic growth. We will consider log-normal noise with shape parameter σ_g . We start with an example in which the parameters are $r = 2$, $K = 8$, $C = 5$ and $\sigma_g = 0.1$.

As a low-dimensional system completely described by three parameters, this scenario should if anything be favorable to a parametric-based approach. This model contains an Allee effect, or tipping point, below which the population is not self-sustaining and shrinks to zero (Courchamp, Berec, and Gascoigne 2008).

Both parametric and nonparametric approaches will require training data on which to base their model of the process. We generate the training data under the model described in Eq 1 for $T_{\text{obs}} = 40$ time steps, under a known but not necessarily optimal sequence of harvest intensities, h_t . For simplicity we imagine a fishery that started from zero harvest pressure and has been gradually increasing the harvest.

Should we include any empirical examples?

The optimal control problem in fisheries management

In our example, we focus on the problem in which a manager must set the harvest level for a marine fishery each year to maximize the net present value of the resource, given an estimated stock size from the year before. Rich data and global concerns have made marine fisheries the crucible for much of the founding work (Gordon 1954; Reed 1979; May et al. 1979; Ludwig and Walters 1982) in managing ecosystems under uncertainty. Global trends (Worm et al. 2006) and controversy (Hilborn 2007; Worm et al. 2009) have made understanding these challenges all the more pressing.

To permit comparisons against a theoretical optimum we will consider data on the stock dynamics simulated

from a simple parametric model in which recruitment of the fish stock X_{t+1} in the following year is a stochastic process governed by a function f of the current stock X_t , selected harvest policy h_t , and noise process Z ,

$$X_{t+1} = Z_t f(X_t, h_t)$$

Given parameters for the function f and probability distribution Z , along with a given economic model determining the price/profit $\Pi(X_t, h_t)$ realized in a given year given a choice of harvest h_t and observed stock X_t . This problem can be solved exactly for discretized values of stock X and policy h using stochastic dynamic programming (SDP) (Mangel 1985). Problems of this sort underpin much marine fisheries management today.

A crux of this approach is correctly specifying the functional form of f , along with its parameters. The standard approach uses one of a handful of common parametric models representing the stock-recruitment relationship, usually after estimating the model parameters from any available existing data. Uncertainty in the parameter estimates can be estimated and integrated over to determine the optimal policy under uncertainty (Mangel 1985; Schapaugh and Tyre 2013). Uncertainty in the model structure itself can only be addressed in this approach by hypothesizing alternative model structures, and then performing some model choice or model averaging (B. K. Williams 2001; Athanassoglou and Xepapadeas 2012). For simplicity of the comparison we consider only the case of a structurally correct and structurally incorrect model, and estimate Bayesian posterior distributions of the parameters. We will also compare to a simpler maximum likelihood estimate of the parameters that ignores the issue of parameter uncertainty.

We compare these parametric approaches to our alternative that uses a Gaussian Process (GP) in place of a given stock recruitment curve. We briefly revisit a description of the Gaussian Process approach that has appeared in recent literature before describing how we take the important step of adapting the GP to the decision theory problem that can still be solved using SDP.

The Non-parametric Bayesian alternative for stock-recruitment curves

The use of Gaussian process (GP) regression (or “kreging” in the geospatial literature) to formulate a predictive model is relatively new in the context of modeling dynamical systems (Kocijan et al. 2005) and introduced in the ecological modeling and fisheries management by Munch et al. (2005). An accessible and thorough introduction to the formulation and use of GPs can be found in Rasmussen and Williams (2006).

The essence of the GP approach can be captured in the following thought experiment: An exhaustive parametric approach to the challenge of structural uncertainty might proceed by writing down all possible functional forms for the underlying dynamical system with all possible parameter values for each form, and then consider searching over this huge space to select the most likely model and parameters; or using a Bayesian approach, assign priors to each of these possible models and infer the posterior distribution of possible models. The GP approach can be thought of as a computationally efficient approximation to this approach. GPs represent a large class of

models that can be thought of as capturing or reasonably approximating the set of models in this collection. By modeling at the level of the process, rather than the level of parametric equation, we can more concisely capture the possible behavior of these curves. In place of a parametric model of the dynamical system, the GP approach postulates a prior distribution of (n-dimensional) curves that can be thought of as approximations to a range of possible (parametric) models that might describe the data. The GP allows us to consider a set of possible curves simultaneously.

Background on Gaussian Process inference

Once the posterior Gaussian process (GP) has been estimated (e.g. see Munch et al. 2005), it is necessary to adapt it in place of the parametric equation for the stochastic dynamic programming (SDP) solution (see Mangel and Clark 1988 for a detailed description of parametric SDP methods) to the optimal policy. The essence of the idea is straight forward – we will use the estimated GP in place of the parametric growth function to determine the stochastic transition matrix on which the SDP calculations are based. The SDP is solved in a discretized state space – both the continuously valued population densities X and harvest quotas h are first mapped to a bounded, discrete grid. (For simplicity we will consider a uniform grid, though for either parametric or GP-based SDP it is often advantageous to use a non-uniform discretization such as a basis function representation, e.g. see (Deisenroth, Rasmussen, and Peters 2009)).

The SDP approach then computes a transition matrix, \mathbf{F} . We demonstrate that calculation is just as straight forward based on the GP as it is in the classical context using the parametric model. The i, j of the transition matrix F entry gives the probability of transitioning into state x_i given that the system is in state x_j in the previous time-step. To generate the transition matrix based on the posterior GP, we need only the expected values at each grid point and the corresponding variances (the diagonal of the covariance matrix), as shown in Figure 1. Given the mean of the GP posterior at each grid-point as the vector E and variance at that point as vector V , the probability of transitioning from state x_i to state x_j is

$$\mathcal{N}\left(x_j | \mu = E_i, \sigma = \sqrt{V_i}\right)$$

where \mathcal{N} is the Normal density at x_j with mean μ and variance σ^2 . Strictly speaking, the transition probability should be calculated by integrating the normal density over the bin of width Δ centered at x_j . For a sufficiently fine grid that $f(x_j) \approx f(x_j + \Delta)$, it is sufficient to calculate the density at x_j and then row-normalize the transition matrix. The process can then be repeated for each possible discrete value of our control variable, (harvest h).

Pseudocode for the determining the transition matrix from the GP

```
for(h in h_grid)
  F_h = for(x_j in grid)
```

```

for(i in 1:N)
  dnorm(x_j, mu[i]-h, V[i])

```

Using the discrete transition matrix we may write down the Bellman recursion defining the stochastic dynamic programming iteration:

$$V_t(x_t) = \max_h \mathbf{E}(h_t + \delta V_{t+1}(Z_{t+1}f(x_t - h_t))) \quad (4)$$

where $V(x_t)$ is the value of being at state x at time t , h is control (harvest level) chosen. Numerically, the maximization is accomplished as follows. Consider the set of possible control values to be the discrete values corresponding the grid of stock sizes. Then for each h_t there is a corresponding transition matrix \mathbf{F}_h determined as described above but with mean $\mu = x_j - h_t$. Let \vec{V}_t be the vector whose i th element corresponds to the value of having stock x_i at time t . Then let Π_h be the vector whose i th element indicates the profit from harvesting at intensity h_t given a population x_i (e.g. $\max(x_i, h_t)$ since one cannot harvest more fish then the current population size). Then the Bellman recursion can be given in matrix form as

$$V_t = \max_h (\Pi_{h_t} + \delta \mathbf{F}_h V_{t+1})$$

where the sum is element by element and the expectation is computed by the matrix multiplication $\mathbf{F}V_{t+1}$.

Pseudocode for the Bellman iteration

```

V1 <- sapply(1:length(h_grid), function(h){
  delta * F[[h]] %*% V + profit(x_grid, h_grid[h])
})
# find havest, h that gives the maximum value
out <- sapply(1:gridsize, function(j){
  value <- max(V1[j,], na.rm = T) # each col is a diff h, max over these
  index <- which.max(V1[j,]) # store index so we can recover h's
  c(value, index) # returns both profit value & index of optimal h.
})
# Sets V[t+1] = max_h V[t] at each possible state value, x
V <- out[1,] # The new value-to-go
D[,OptTime-time+1] <- out[2,] # The index positions

```

This completes the algorithm adapting the GP to the sequential decision-making problem through SDP, which we believe has not yet been demonstrated in the literature. We provide an R package implementation of this, along with the Gaussian process inference, in the supplemental materials.

Estimating parametric models

We estimate posterior distributions for two parametric models: one using the structurally correct model as given in Eq (1), which we refer to as the “Parametric Bayes” model, and another using the familiar Ricker model, using a Gibbs sampler as described (with source code) in the appendix). In addition we estimate the parameters of the structurally correct model by maximum likelihood.

The posterior distribution for the hyper-parameters of the Gaussian process model are estimated by Metropolis-Hastings algorithm, again with details and code provided in the Appendix. Rasmussen and Williams (2006) provides an excellent general introduction to Gaussian Processes and Munch et al. (2005) first discusses their application in this context.

Results

```
require(MASS)
step_ahead <- function(x, f, p){
  h = 0
  x_predict <- sapply(x, f, h, p)
  n <- length(x_predict) - 1
  y <- c(x[1], x_predict[1:n])
  y
}

step_ahead_posteriors <- function(x){
  gp_f_at_obs <- gp_predict(gp, x, burnin=1e4, thin=300)
  df_post <- melt(lapply(sample(100),
    function(i){
      data.frame(time = 1:length(x), stock = x,
        GP = mvrnorm(1, gp_f_at_obs$Ef_posterior[,i], gp_f_at_obs$Cf_posterior[[i]]),
        True = step_ahead(x,f,p),
        MLE = step_ahead(x,f,est$p),
        Allen = step_ahead(x, allen_f, pardist[i,]),
        Ricker = step_ahead(x, ricker_f, ricker_pardist[i,]),
        Myers = step_ahead(x, myers_f, myers_pardist[i,]))
    }, id=c("time", "stock")))
}

df_post <- step_ahead_posteriors(x)
```



```
ggplot(df_post) + geom_point(aes(time, stock)) +
  geom_line(aes(time, value, col=variable, group=interaction(L1,variable)), alpha=.1) +
  scale_colour_manual(values=colorkey, guide = guide_legend(override.aes = list(alpha = 1)))
```




Figure 2: plot of chunk Figureb

Figure 1 shows the mean inferred state space dynamics of each model relative to the true model used to generate the data, predicting the relationship between observed stock size (x-axis) to the stock size after recruitment the following year. All models except the MLE model estimate a distribution around the means shown here, and all models estimate a level of process noise, which is independent of the state value (x). Note that in contrast to the other models shown, the mean Gaussian process corresponds to a distribution of curves - as indicated by the gray band - which itself has a mean shown in black. Note that this mean GP is thus more certain of the dynamics in the region where data is available than where it is not.

While it would be natural (and straight forward) to condition the GP on passing through the origin (0,0) (see appendix), the estimate shown here is based only on the observed data. The observed data from which each model is estimated is also shown. The observations come from only a limited region of state space corresponding to unharvested or weakly harvested system. No observations occur at the theoretical optimum or near the tipping point.

```
policies <- melt(data.frame(stock=x_grid, sapply(OPT, function(x) x_grid[x])), id="stock")
```

```
names(policies) <- c("stock", "method", "value")

ggplot(policies, aes(stock, stock - value, color=method)) +
  geom_line(lwd=1.2, alpha=0.8) + xlab("stock size") + ylab("escapement") +
  scale_colour_manual(values=colorkey)
```

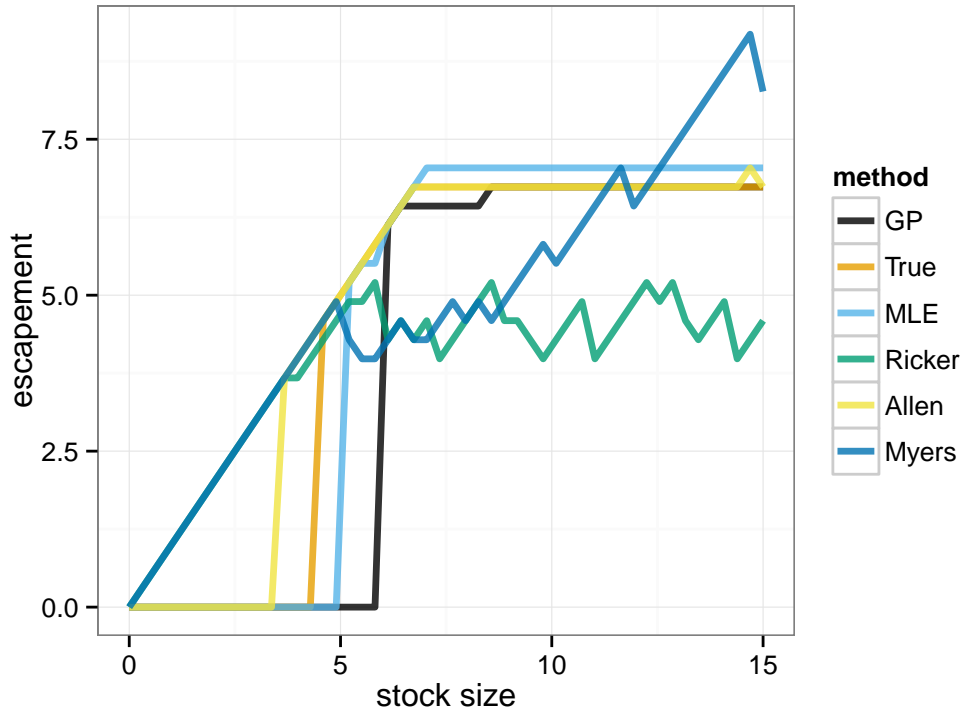


Figure 3: The steady-state optimal policy (infinite boundary) calculated under each model. Policies are shown in terms of target escapement, S_t , as under models such as this a constant escapement policy is expected to be optimal (Reed 1979).

The resulting optimal management strategy based on each of the inferred models is shown in Figure 2, against the optimal strategy given the true underlying dynamics. Policies are shown in terms of target escapement, S_t . Under models such as this a constant escapement policy is expected to be optimal (Reed 1979), whereby population levels below a certain size S are unharvested, while above that size the harvest strategy aims to return the population to S , resulting in the hockey-stick shaped policies shown.

```
ggplot(sims_data) +
  geom_line(aes(time, fishstock, group=interaction(reps,method), color=method), alpha=.1) +
  scale_colour_manual(values=colorkey, guide = guide_legend(override.aes = list(alpha = 1)))
```

The consequences of managing 100 replicate realizations of the simulated fishery under each of the policies estimated is shown in Figure 3. As expected from the policy curves, the structurally correct model under-harvests,




Figure 4: Gaussian process inference outperforms parametric estimates. Shown are 100 replicate simulations of the stock dynamics (eq 1) under the policies derived from each of the estimated models, as well as the policy based on the exact underlying model.

leaving the stock to vary around it's un-fished optimum. The structurally incorrect Ricker model over-harvests the population passed the tipping point consistently, resulting in the immediate crash of the stock and thus derives minimal profits.

The results shown in Figures 1-3 are not unique to the simulated data or models chosen here, but arises across a range of parameter values and simulations as shown in the supplemental figures. The results across this range can most easily be compared by the relative differences in net present value realized by each of the approaches, as shown in Figure 4. The Gaussian Process most consistently realizes a value close to the optimal solution, and importantly avoids ever driving the system across the tipping point, which results in the near-zero value cases in the parametric models.

```
ggplot(actual_over_optimal, aes(value)) + geom_histogram(aes(fill=variable)) +  
  facet_wrap(~variable, scales = "free_y") + guides(legend.position = "none") +  
  xlab("Total profit by replicate") + scale_fill_manual(values=colorkey)
```

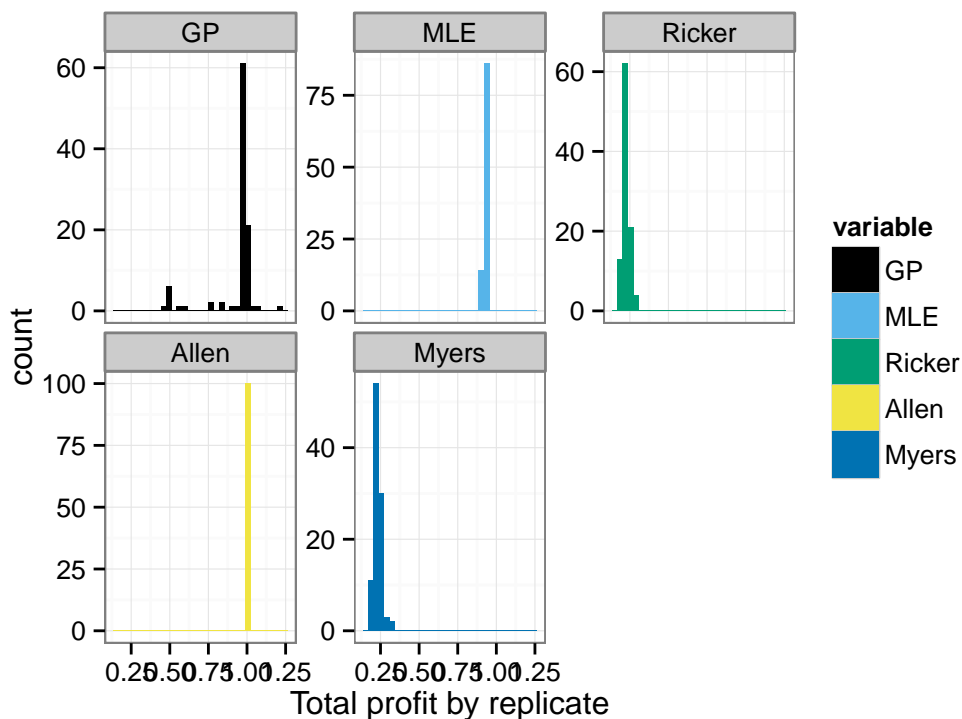


Figure 5: Histograms of the realized net present value of the fishery over a range of simulated data and resulting parameter estimates. For each data set, the three models are estimated as described above. Values plotted are the averages of a given policy over 100 replicate simulations. Details and code provided in the supplement.

```
ggplot(actual_over_optimal, aes(value)) + geom_histogram(aes(fill=variable), binwidth=0.1) +  
  xlab("Total profit by replicate")+ scale_fill_manual(values=colorkey)
```




Figure 6: Histograms of the realized net present value of the fishery over a range of simulated data and resulting parameter estimates. For each data set, the three models are estimated as described above. Values plotted are the averages of a given policy over 100 replicate simulations. Details and code provided in the supplement.

```
ggplot(actual_over_optimal, aes(value, fill=variable, color=variable)) +
  stat_density(aes(y=..density..), position="stack", adjust=3, alpha=.9) +
  xlab("Total profit by replicate")+ scale_fill_manual(values=colorkey)+ scale_color_manual(values=colorkey)
```

Discussion

Non-parametric Bayesian methods have received far too little attention in ecological modeling efforts that are aimed at improved conservation planning and decision making support. Such approaches may be particularly useful when the available data is restricted to a limited area of state-space, which can lead parametric models to underestimate the uncertainty in dynamics at population levels (states) which have not been observed. One reason for the relative absence of nonparametric approaches in the natural resource management context may be the lack of existing approaches for adapting the non-parametric Bayesian models previously proposed (Munch et al. 2005) to a decision-theoretic framework. Adapting a non-parametric approach requires modification of existing methods for decision theory. We have illustrated how this might be done for a classic stochastic dynamic




Figure 7: Histograms of the realized net present value of the fishery over a range of simulated data and resulting parameter estimates. For each data set, the three models are estimated as described above. Values plotted are the averages of a given policy over 100 replicate simulations. Details and code provided in the supplement.

programming problem, opening the door for substantial further research into how these applications might be improved.

While non-parametric Bayesian approaches will not always be preferable to simple mechanistic models, we highlight three aspects of the problem consider here that make these methods particularly valuable. These aspects are common to many conservation decision making problems, which thus merit greater use of non-parametric approaches that can best take advantage of them.

1. Large uncertainty where the data is poor

The parametric models perform worst when they propose a management strategy outside the range of the observed data. The non-parametric Bayesian approach, in contrast, allows a predictive model that expresses a great deal of uncertainty about the probable dynamics outside the observed range, while retaining very good predictive accuracy in the range observed. The management policy dictated by the GP balance this uncertainty against the immediate value of the harvest, and act to stabilize the population dynamics in a region of state space in which the predictions can be reliably reflected by the data.

2. Predictive accuracy where data is good

While expressing larger uncertainty outside the observed data, the GP can also provide a better fit with smaller uncertainty inside the range of the observed data. This arises from the greater flexibility of the Gaussian process, which describes a large family of possible curves.

While in a parametric context this over-fitting would be more worrisome – a high-degree polynomial could fit the data even better – those concerns are driven by the resulting parametric fit outside the data, which may involve wild oscillations unsupported by the data. As we have seen in #1, the GP is less vulnerable to such unjustified predictions outside the data, and is meanwhile free to benefit from the greater fit where the data is available.

Future directions

In this simulated example, the underlying dynamics are truly governed by a simple parametric model, allowing the parametric approaches to be more accurate. Similarly, because the dynamics are one-dimensional dynamics and lead to stable nodes (rather than other attractors such as limit-cycles resulting in oscillations), the training data provides relatively limited information about the dynamics. For these reasons, we anticipate that in higher-dimensional examples characteristic of ecosystem management problems that the machine learning approach will prove even more valuable.

Data complexity. Perhaps too far out of scope... The nonparametric Bayesian approach is also better suited to complex and disparate data. Incorporating various sources of information into mechanistic models can be an immensely difficult due to the increased complexity involved.

Online learning

In our treatment here we have ignored the possibility of learning during the management phase, in which the additional observations of the stock size could potentially improve parameter estimates. While we intend to address this possibility in future work in the context of these non-parametric models, we have not addressed it here for pedagogical reasons. In the context presented here, it is clear that the differences in performance arise from differences in the uncertainty inherent in the model formulations, rather than from differing abilities to learn. Because we consider a threshold system, online learning would not change this generic feature of a lack of data in a certain range of the state space which is better captured by the Gaussian process.

Acknowledgments

This work was partially supported by the Center for Stock Assessment Research, a partnership between the University of California Santa Cruz and the Fisheries Ecology Division, Southwest Fisheries Science Center, Santa Cruz, CA and by NSF grant EF-0924195 to MM.

Appendix

The appendices have not yet been assembled. Meanwhile, code to repeat the analyses, along with a complete log of all research conducted on this project, can be found at: <https://github.com/cboettig/nonparametric-bayes>

Markov Chain Monte Carlo Analysis

```
gp_assessment_plots[[1]]
```

```
gp_assessment_plots[[2]]
```

```
plot_allen_traces
```

```
plot_allen_posteriors
```

```
plot_ricker_traces
```

```
plot_ricker_posteriors
```

```
plot_myers_traces
```

```
plot_myers_posteriors
```





Figure 8: plot of chunk appendixplots

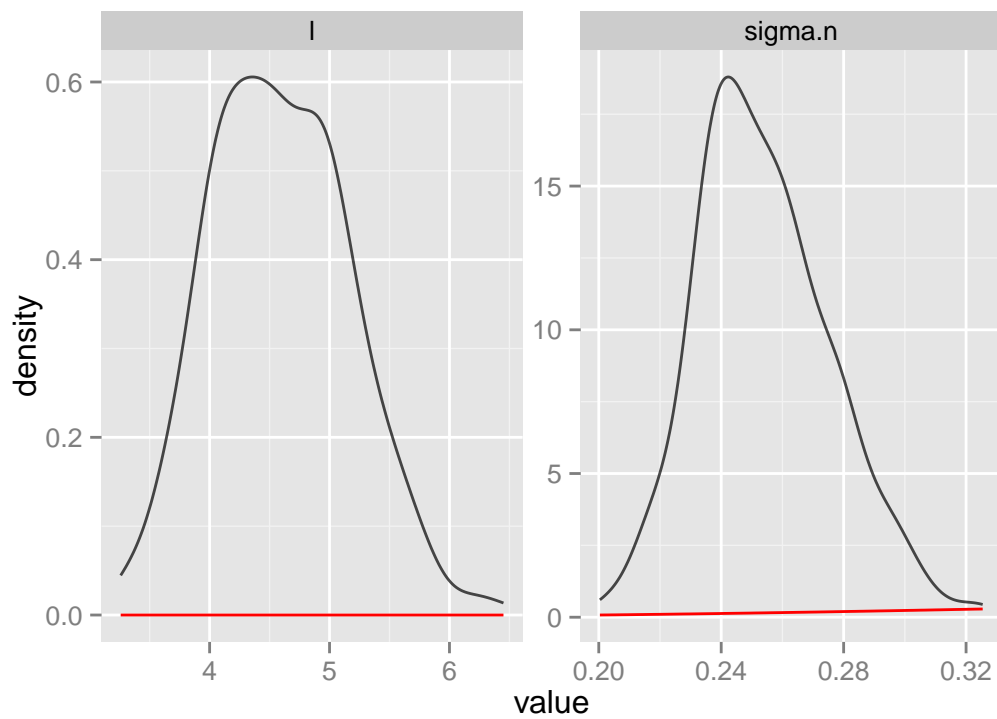


Figure 9: plot of chunk appendixplots




Figure 10: plot of chunk appendixplots

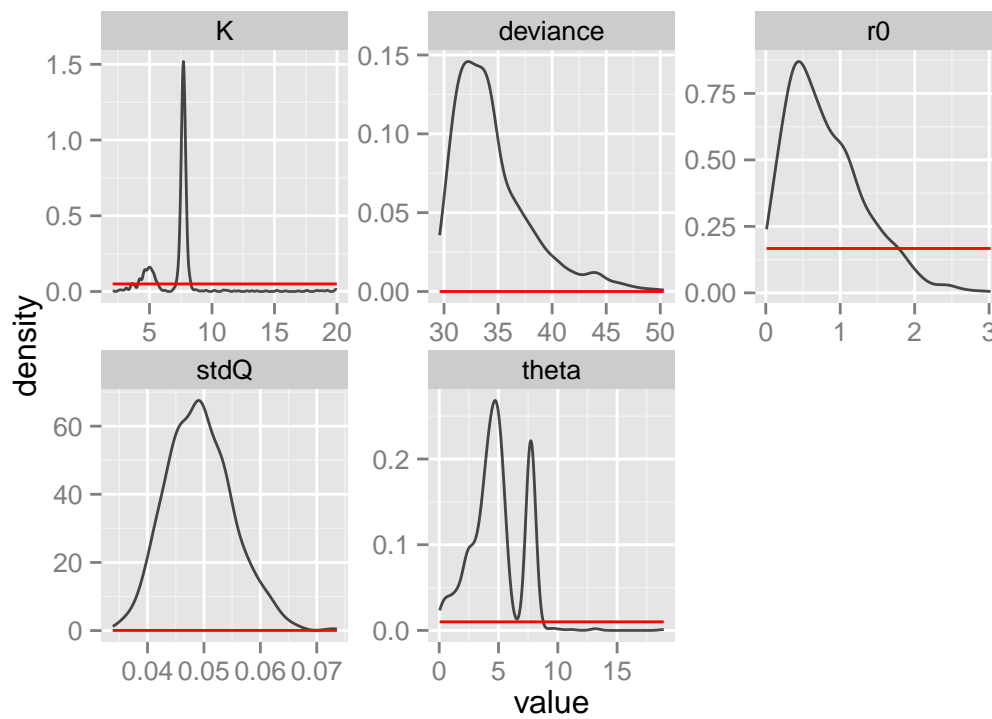


Figure 11: plot of chunk appendixplots




Figure 12: plot of chunk appendixplots




Figure 13: plot of chunk appendixplots




Figure 14: plot of chunk appendixplots

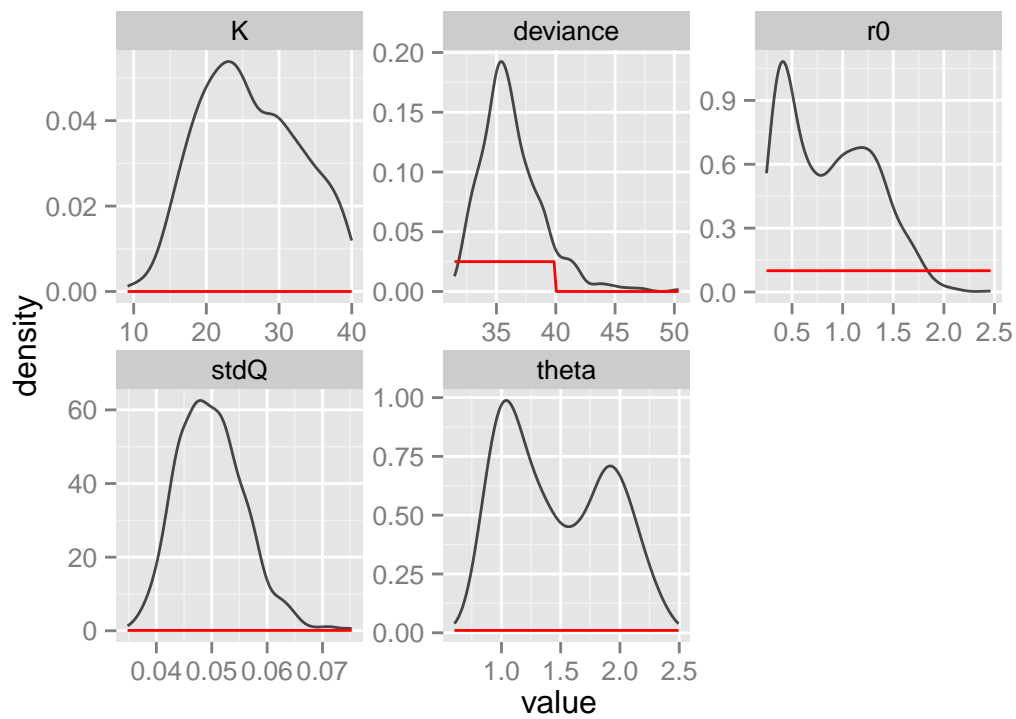


Figure 15: plot of chunk appendixplots

References

- Allen, Linda J. S., Jesse F. Fagan, Göran Högnäs, and Henrik Fagerholm. 2005. "Population extinction in discrete-time stochastic population models with an Allee effect." *Journal of Difference Equations and Applications* 11 (4-5) (apr): 273–293. doi:10.1080/10236190412331335373. <http://www.tandfonline.com/doi/abs/10.1080/10236190412331335373>.
- Athanassoglou, Stergios, and Anastasios Xepapadeas. 2012. "Pollution control with uncertain stock dynamics: When, and how, to be precautionous." *Journal of Environmental Economics and Management* 63 (3) (may): 304–320. doi:10.1016/j.jeem.2011.11.001. <http://linkinghub.elsevier.com/retrieve/pii/S0095069611001409>.
- Barnosky, Anthony D., Elizabeth a Hadly, Jordi Bascompte, Eric L. Berlow, James H. Brown, Mikael Fortelius, Wayne M. Getz, et al. 2012. "Approaching a state shift in Earth's biosphere." *Nature* 486 (7401) (jun): 52–58. doi:10.1038/nature11018. <http://www.nature.com/doifinder/10.1038/nature11018>.
- Bestelmeyer, Brandon T., Michael C. Duniway, Darren K. James, Laura M. Burkett, and Kris M. Havstad. 2012. "A test of critical thresholds and their indicators in a desertification-prone ecosystem: more resilience than we thought." Ed. Katharine Suding. *Ecology Letters* (dec). doi:10.1111/ele.12045. <http://doi.wiley.com/10.1111/ele.12045>.
- Boettiger, Carl, and Alan Hastings. 2013. "Tipping points: From patterns to predictions." *Nature* 493 (7431) (jan): 157–158. doi:10.1038/493157a. <http://www.nature.com/doifinder/10.1038/493157a>.
- Brozović, Nicholas, and Wolfram Schlenker. 2011. "Optimal management of an ecosystem with an unknown threshold." *Ecological Economics* (jan): 1–14. doi:10.1016/j.ecolecon.2010.10.001. <http://linkinghub.elsevier.com/retrieve/pii/S0921800910004167>.
- Courchamp, Franck, Ludek Berec, and Joanna Gascoigne. 2008. *Allee Effects in Ecology and Conservation*. Oxford University Press, USA. <http://www.amazon.com/Effects-Ecology-Conservation-Franck-Courchamp/dp/0198570309>.
- Cressie, Noel, Catherine a Calder, James S. Clark, Jay M. Ver Hoef, and Christopher K. Wikle. 2009. "Accounting for uncertainty in ecological analysis: the strengths and limitations of hierarchical statistical modeling." *Ecological Applications* 19 (3) (apr): 553–70. <http://www.ncbi.nlm.nih.gov/pubmed/19425416>.
- Deisenroth, Marc Peter, Carl Edward Rasmussen, and Jan Peters. 2009. "Gaussian process dynamic programming." *Neurocomputing* 72 (7-9) (mar): 1508–1524. doi:10.1016/j.neucom.2008.12.019. <http://linkinghub.elsevier.com/retrieve/pii/S0925231209000162>.
- Geritz, Stefan A. H., and Eva Kisdi. 2012. "Mathematical ecology: why mechanistic models?." *Journal of mathematical biology* 65 (6-7) (dec): 1411–5. doi:10.1007/s00285-011-0496-3. <http://www.ncbi.nlm.nih.gov/pubmed/22159789>.

- Gordon, H. S. 1954. "The economic theory of a common-property resource: the fishery." *The Journal of Political Economy* 62 (2): 124–142. <http://www.jstor.org/stable/10.2307/1825571>.
- Hilborn, Ray. 2007. "Reinterpreting the State of Fisheries and their Management." *Ecosystems* 10 (8) (oct): 1362–1369. doi:10.1007/s10021-007-9100-5. <http://www.springerlink.com/index/10.1007/s10021-007-9100-5>.
- Hughes, Terry P., Cristina Linares, Vasilis Dakos, Ingrid a van de Leemput, and Egbert H. van Nes. 2013. "Living dangerously on borrowed time during slow, unrecognized regime shifts." *Trends in ecology & evolution* 28 (3) (mar): 149–55. doi:10.1016/j.tree.2012.08.022. <http://www.ncbi.nlm.nih.gov/pubmed/22995893>.
- Kocijan, Juš, Agathe Girard, Blaž Banko, and Roderick Murray-Smith. 2005. "Dynamic systems identification with Gaussian processes." *Mathematical and Computer Modelling of Dynamical Systems* 11 (4) (dec): 411–424. doi:10.1080/13873950500068567. <http://www.tandfonline.com/doi/abs/10.1080/13873950500068567>.
- Ludwig, Donald, and Carl J. Walters. 1982. "Optimal harvesting with imprecise parameter estimates." *Ecological Modelling* 14 (3-4) (jan): 273–292. doi:10.1016/0304-3800(82)90023-0. <http://linkinghub.elsevier.com/retrieve/pii/0304380082900230>.
- Mangel, Marc. 1985. "Decision and control in uncertain resource systems:" 255. <http://www.amazon.com/Decision-uncertain-resource-Mathematics-Engineering/dp/0124687202> <http://dl.acm.org/citation.cfm?id=537497>.
- Mangel, Marc, and Colin W. Clark. 1988. *Dynamic Modeling in Behavioral Ecology*. Ed. John Krebs and Tim Clutton-Brock. Princeton: Princeton University Press.
- May, Robert M., John R. Beddington, Colin W. Clark, Sidney J. Holt, and R. M. Laws. 1979. "Management of multispecies fisheries." *Science (New York, N.Y.)* 205 (4403) (jul): 267–77. doi:10.1126/science.205.4403.267. <http://www.ncbi.nlm.nih.gov/pubmed/17747032>.
- Munch, Stephan B., Melissa L. Snover, George M. Watters, and Marc Mangel. 2005. "A unified treatment of top-down and bottom-up control of reproduction in populations." *Ecology Letters* 8 (7) (may): 691–695. doi:10.1111/j.1461-0248.2005.00766.x. <http://doi.wiley.com/10.1111/j.1461-0248.2005.00766.x>.
- Rasmussen, Carl Edward, and C. K. I. Williams. 2006. *Gaussian Processes for Machine Learning*. Ed. Thomas Dietterich. Boston: MIT Press,. www.GaussianProcess.org/gpml.
- Reed, William J. 1979. "Optimal escapement levels in stochastic and deterministic harvesting models." *Journal of Environmental Economics and Management* 6 (4) (dec): 350–363. doi:10.1016/0095-0696(79)90014-7. <http://www.sciencedirect.com/science/article/pii/0095069679900147> <http://linkinghub.elsevier.com/retrieve/pii/0095069679900147>.
- Schapaugh, Adam W., and Andrew J. Tyre. 2013. "Accounting for parametric uncertainty in Markov decision processes." *Ecological Modelling* 254 (apr): 15–21. doi:10.1016/j.ecolmodel.2013.01.003. <http://linkinghub.elsevier.com/retrieve/pii/S0304380013000306>.
- Scheffer, Marten, Jordi Bascompte, William A. Brock, Victor Brovkin, Stephen R. Carpenter, Vasilis Dakos,

- Hermann Held, Egbert H. van Nes, Max Rietkerk, and George Sugihara. 2009. “Early-warning signals for critical transitions.” *Nature* 461 (7260): 53–9. <http://www.ncbi.nlm.nih.gov/pubmed/19727193>.
- Scheffer, Marten, Stephen R. Carpenter, J. A. Foley, C. Folke, and B. Walker. 2001. “Catastrophic shifts in ecosystems.” *Nature* 413 (6856) (oct): 591–6. doi:10.1038/35098000. <http://www.ncbi.nlm.nih.gov/pubmed/11595939>.
- Williams, Byron K. 2001. “Uncertainty , learning , and the optimal management of wildlife.” *Environmental and Ecological Statistics* 8: 269–288. doi:10.1023/A:1011395725123.
- Worm, Boris, Edward B. Barbier, Nicola Beaumont, J. Emmett Duffy, Carl Folke, Benjamin S. Halpern, Jeremy B. C. Jackson, et al. 2006. “Impacts of biodiversity loss on ocean ecosystem services.” *Science (New York, N.Y.)* 314 (5800) (nov): 787–90. doi:10.1126/science.1132294. <http://www.ncbi.nlm.nih.gov/pubmed/17082450>.
- Worm, Boris, Ray Hilborn, Julia K. Baum, Trevor A. Branch, Jeremy S. Collie, Christopher Costello, Michael J. Fogarty, et al. 2009. “Rebuilding global fisheries.” *Science (New York, N.Y.)* 325 (5940) (jul): 578–85. doi:10.1126/science.1173146. <http://www.ncbi.nlm.nih.gov/pubmed/19644114>.