

Avoiding tipping points in fisheries management through Gaussian Process Dynamic Programming

Carl Boettiger^{*,a}, Marc Mangel^a, Stephan Munch^c

^a*Center for Stock Assessment Research, Department of Applied Math and Statistics, University of California, Mail Stop SOE-2, Santa Cruz, CA 95064, USA*

^b*Center for Stock Assessment Research, Department of Applied Math and Statistics, University of California, Mail Stop SOE-2, Santa Cruz, CA 95064, USA and Department of Biology, University of Bergen, Bergen, Norway 9020*

^c*Southwest Fisheries Science Center, National Oceanic and Atmospheric Administration, 110 Shaffer Road, Santa Cruz, CA 95060, USA*

Abstract

Model uncertainty and limited data are fundamental challenges to robust management of human intervention in a natural system. These challenges are acutely highlighted by concerns that many ecological systems may contain tipping points, such as Allee population sizes. Before a collapse, we do not know where the tipping points lie, if they exist at all. Hence, we know neither a complete model of the system dynamics nor do we have access to data in some large region of state-space where such a tipping point might exist. We illustrate how a Bayesian Non-Parametric (BNP) approach using a Gaussian Process (GP) prior provides a flexible representation of this inherent uncertainty. We embed GPs in a Stochastic Dynamic Programming (SDP) framework in order to make robust management predictions with both model uncertainty and limited data. We use simulations to evaluate this approach as compared with the standard approach of using model selection to choose from a set of candidate models. We find that model selection erroneously favors models without tipping points – leading to harvest policies that guarantee extinction. The GPDP performs nearly as well as the true model and significantly outperforms standard approaches. We illustrate this using examples of simulated single-species dynamics, where the standard model selection approach should be most effective, and find that it still fails to account for uncertainty appropriately and leads to population crashes, while management based on the GPDP does not, since it does not underestimate the uncertainty outside of the observed data.

Key words: Bayesian, Structural Uncertainty, Nonparametric, Optimal Control, Decision Theory, Gaussian Processes, Fisheries Management,

^{*}Corresponding author

Email address: cboettig(at)gmail.com (Carl Boettiger)

14 Introduction

15 Decision making under uncertainty is a ubiquitous challenge in the management of human
16 intervention in natural resources and conservation. Decision-theoretic approaches provide a
17 framework to determine the best sequence of actions in face of uncertainty, but only when
18 that uncertainty can be meaningfully quantified (Fischer et al. 2009). Over the last four
19 decades (beginning with Clark (1976), Clark (2009) and Walters and Hilborn (1978)) dynamic
20 optimization methods, particularly Stochastic Dynamic Programming (SDP), have become
21 increasingly important as a means of understanding how to manage human intervention into
22 natural systems. Simultaneously, there has been increasing recognition of the importance of
23 multiple steady states or ‘tipping points’ (Scheffer et al. 2001, 2009, Polasky et al. 2011) in
24 ecological systems.

25 We develop a novel approach to address these concerns in the context of fisheries; although
26 the challenges and methods are germane to other problems of conservation or natural resource
27 exploitation. Economic value and ecological concern have made marine fisheries the crucible for
28 much of the founding work on management under uncertainty (Gordon 1954, Clark 1976, 2009,
29 May et al. 1979, Reed 1979, Ludwig and Walters 1982).

30 Even if we know the proper deterministic description of the biological system, there is intrinsic
31 stochasticity in biological dynamics, measurements, and implementation of policy (*e.g.* Reed
32 1979, Clark and Kirkwood 1986, Roughgarden and Smith 1996, Sethi et al. 2005). We may also
33 lack knowledge about the parameters of the biological dynamics (parametric uncertainty, *e.g.*
34 Ludwig and Walters 1982, Hilborn and Mangel 1997, McAllister 1998, Schapaugh and Tyre 2013),
35 or even not know which model is proper description of the system (structural uncertainty, *e.g.*
36 Williams 2001, Cressie et al. 2009, Athanassoglou and Xepapadeas 2012). Of these, the latter
37 is generally the hardest to quantify. Typical approaches confront the data with a collection of
38 models, assuming that the true dynamics (or reasonable approximation) is among the collection
39 and then use model choice or model averaging to arrive at a conclusion (Williams 2001, Cressie
40 et al. 2009, Athanassoglou and Xepapadeas 2012). Even setting aside other concerns (see Cressie
41 et al. (2009)), these approaches are unable to describe uncertainty outside the observed data
42 range.

43 Structural uncertainty is particularly insidious when we try to predict outside of the range of
44 observed data (Mangel et al. 2001) because we are extrapolating into unknown regions. In
45 management applications, this extrapolation uncertainty is particularly important since (a)
46 management involves considering actions that may move the system outside the range of observed
47 behavior, and (b) the decision tools (optimal control theory, SDP) rely on both reasonable
48 estimates of the expected outcomes and on the weights given to those outcomes (*e.g.* Weitzman
49 2013). Thus characterizing uncertainty is as important as characterizing the expected outcome.
50 Tipping points in ecological dynamics (Scheffer et al. 2001, Polasky et al. 2011) highlight this
51 problem because precise models are not available and data are limited such as around high stock
52 levels or an otherwise desirable state. With perfect information, one would know just how far a
53 system could be pushed before crossing the tipping point, and management would be simple.
54 But we face imperfect models and limited data and, with tipping points,
55 even small errors can have very large consequences, as we shall illustrate later. Because
56 intervention may be too late once a tipping point has been crossed (but see Hughes et al. (2013)),
57 management is often concerned with avoiding tipping points before any data about them are
58 available.

59 The dual concerns of model uncertainty and incomplete data create a substantial challenge
60 to existing decision-theoretic approaches (Brozović and Schlenker 2011). We illustrate how
61 Stochastic Dynamic Programming (SDP) (Mangel and Clark 1988, Marescot et al. 2013) can be
62 implemented using a Bayesian Non-Parametric (BNP) model of population dynamics (Munch et
63 al. 2005a). The BNP method has two distinct advantages. First, using a BNP model sidesteps
64 the need for an accurate model-based description of the system dynamics. Second, a BNP model
65 better reflects uncertainty when extrapolating beyond the observed data. This is crucial to
66 providing robust decision-making when the correct model is not known (as is almost always
67 true). [We use *robust* to characterize approaches that provide nearly optimal solutions without
68 being sensitive to the choice of the (unknown) underlying model.]

69 This paper is the first ecological application of the SDP without an *a priori* model of the underlying
70 dynamics. Unlike parametric approaches that can only reflect uncertainty in parameter estimates,
71 the BNP method provides a broader representation of uncertainty, including uncertainty beyond

the observed data. We will show that Gaussian Process Dynamic Programming (GPDP) allows us to find robust management solutions in face of limited data without knowing the correct model structure.

For comparisons, we consider the performance of management based on GPDP against management policies derived under several alternative parametric models (Reed 1979, Ludwig and Walters 1982, Mangel and Clark 1988). Rather than compare models in terms of best fit to data, we compare model performance in the concrete terms of the decision-maker’s objectives.

Approach and Methods

We first describe the requirements of dynamic optimization for the management of human intervention in natural resource systems. After that we describe three parametric models for population dynamics and the Gaussian Process (GP)¹ description of population dynamics.

Requirements of dynamic optimization

Dynamic optimization requires characterizing the dynamics of a state variable (or variables), a control action, and a value function. For simplicity, we consider only a single state variable, which allows us to focus on one-dimensional models of population dynamics for the comparison of parametric models and the GP. This is a best-case scenario for the parametric models because we simulate underlying dynamics from one of the three parametric models, whereas in the natural world we never know the “true” model. In addition, by choosing one-dimensional models with just a few parameters, we limit the chance that poor performance will be due to inability to estimate parameters accurately, something that becomes a more severe problem for higher-dimensional parametric models. Finally, the parametric models we consider are commonly used in modeling stock-recruitment dynamics or to model sudden transitions between alternative stable states.

¹We abbreviate Gaussian Process as GP, which refers to the statistical model we use to approximate the population dynamics, and we use the term Gaussian Process Dynamic Programming [GPDP], to refer to the use of a GP as the underlying process model when solving a Dynamic Programming equation. Hence we will refer to the models as: GP, Ricker, Allen, etc, and the novel method we put forward here as GPDP.

95 We let $X(t)$ denote the size (numbers or biomass) of the focal population at time t and assume
 96 that in the absence of take its dynamics are:

$$X(t+1) = Z(t)f(X(t), \mathbf{p}) \quad (1)$$

97 Where $Z(t)$ is log-normally distributed process stochasticity (Reed 1979) and \mathbf{p} is a vector of
 98 parameters to be estimated from the data. We describe the three choices for $f(X(t), \mathbf{p})$ in the
 99 next section.

100 The control action is a harvest or take, $h(t)$, measured in the same units as X , at time t . Thus,
 101 in the presence of take, the population size on the right hand side of Eqn 1 is replaced by
 102 $S(t) = X(t) - h(t)$.

103 To construct the value function, we consider a return when $X(t) = x(t)$ and harvest $h(t) = h$
 104 denoted as the reward, $R(x(t), h)$. For example, if the return is the harvest at time t , then
 105 $R(x(t), h(t)) = \min(x(t), h(t))$. We assume that future harvests are discounted relative to
 106 current ones at a constant rate of discount δ and ask for the harvest policy that maximizes total
 107 discounted harvest between the current time t and a final time T . That is, we seek to maximize
 108 over choices of harvest $E[\sum_{t=0}^T R(X(t), h(t), t)\delta^t]$, where the state dynamics are given by Eqn 1
 109 and E denotes the expectation over future population states.

110 In order to find that policy, we introduce the value function $V(x(t), t)$ representing the total
 111 discounted catch from time t onwards given that $X(t) = x(t)$. This value function satisfies an
 112 equation of SDP (Mangel and Clark 1988, Clark and Mangel 2000, Clark 2009 Mangel 2014),

$$V(x(t), t) = \max_{h(t)} \{R(h(t), x(t)) + \delta \cdot E_{X(t+1)} [V(X(t+1), t+1)|x(t), h(t)]\} \quad (2)$$

113 where expectation is taken over all possible values of the next state, $X(t+1)$, and maximized
 114 over all possible choices of harvest, $h(t)$. That is, at time t , when population size is $x(t)$ and
 115 harvest $h(t)$ is applied, the immediate return is $R(h(t), x(t))$. When the sole source of uncertainty
 116 is the process stochasticity term, Z , and thus the expectation in Eqn 2 is equivalent to taking
 117 expectations over $Z(t)$. That is

$$E_{X(t+1)} [V(X(t+1), t+1)|x(t), h(t)] = E_{Z(t)} [V(Z(t)f(x(t) - h(t))|\mathbf{p}), t+1|x(t), h(t)] \quad (3)$$

118 where the population size after the take is $x(t) - h(t)$, which is then translated into $X(t+1)$ by
 119 Eqn 1 (that is, we replace $X(t+1)$ by $Z(t)f(x(t) - h(t)|\mathbf{p})$).

120 When the parameters governing the dynamics are also uncertain (the case of parametric un-
 121 certainty), the expectation in Eqn 2 involves averaging over the posterior distribution for the
 122 parameters. That is,

$$E_{X(t+1)} [V(X(t+1), t+1)|x(t), h(t)] = E_{\mathbf{p}|\text{data}} \{E_{Z(t)|\mathbf{p}, \text{data}} [V(Z(t)f(x(t) - h(t))|\mathbf{p}), t]\} \quad (4)$$

123 When the underlying population dynamics are unknown (the case of structural uncertainty), the
 124 function f itself is uncertain and the expectation for the next state includes uncertainty in f as
 125 well. That is

$$E_{X(t+1)} [V(X(t+1), t+1)|x(t), h(t)] = E_{\mathbf{p}|\text{data}} \{E_{f, Z(t)|\mathbf{p}, \text{data}} [V(Z(t)f(x(t) - h(t))|\mathbf{p}), t]\} \quad (5)$$

126 We consider the finite time problem with $T = 1000$, which we solve using the standard value
 127 iteration algorithm (see Mangel and Clark 1988, Clark and Mangel 2000).

128 *Parametric models*

129 We consider three candidate parametric models for the population dynamics: The Ricker model,
 130 the Allen model (Allen and Tanner 2005), and the Myers model (Myers et al. 1995), Eqns (6)-(8).
 131 In all three, we let K denote the carrying capacity and r the maximum per capita growth rate.
 132 The Ricker model has two parameters and the right hand side of Eqn 1 is

$$f(S(t)|r, K) = S(t)e^{r(1 - \frac{S(t)}{K})} \quad (6)$$

133 The Allen model has three parameters

$$f(S(t)|r, K, X_c) = S(t)e^{r(1-\frac{S(t)}{K})(S(t)-X_c)} \quad (7)$$

134 where X_c denotes the location of the unstable steady state (i.e., the tipping point).

135 The Myers model also has three parameters

$$f(S(t)|r, K, \theta) = \frac{rS(t)^\theta}{1 + \frac{S(t)^\theta}{K}} \quad (8)$$

136 where $\theta = 1$ corresponds to Beverton-Holt dynamics and $\theta > 2$ leads to Allee effects and multiple
137 stable states.

138 The Ricker model does not lead to multiple steady states. The Allen model resembles the Ricker
139 dynamics with an added Allee effect parameter (Courchamp et al. 2008), below which the
140 population cannot persist. The Myers model also has three parameters and contains an Allee
141 threshold, but unlike the Ricker model saturates at high population size. The multiplicative
142 log-normal stochasticity in Eqn 1 introduces one additional parameter σ that must be estimated.
143 Because of our interest in management performance in the presence of tipping points, all of our
144 simulations are based on the Allen model. The Allen model is thus the state of nature and is
145 expected to provide the best-case scenario. The Ricker model is a reasonable approximation of
146 these dynamics far from the Allee threshold (but lacks threshold dynamics), while the Myers
147 model shares the essential feature of a threshold but differs in structure from the Allen model.
148 Throughout, we refer to the “True” model when the underlying parameters *are known without*
149 *error*, and refer to the “Allen” model when these parameters have been estimated from the
150 sample data.

151 We consider a period of 40 in which data are obtained to estimate the parameters or the GP.
152 This is long enough that the estimates do not depend on the particular realization, and longer
153 times are not likely to provide substantial improvement. Each of the models is fit to the same
154 data (Figure 1).

155 We inferred posterior distributions for the parameters of each model by Gibbs sampling (Gelman
156 et al. (2003) implemented in R (R Core Team 2013) using `jags`, (Su and Masanao Yajima 2013)).

157 We choose uniform priors for all parameters (See appendix Tables S1-S3; R code provided).
158 We show one-step-ahead predictions of these model fits in Figure 1. We tested each chain for
159 Gelman-Rubin convergence and results were robust to longer runs. For each simulation we
160 also applied several commonly used model selection criteria (AIC, BIC, DIC, see Burnham and
161 Anderson (2002)) to identify the best fitting model.

162 Additionally, we compute the maximum likelihood estimate (MLE, as we will refer to this model
163 in the figures) of the parameters for the (structurally correct) Allen model. Comparing this to
164 using the posterior distribution of parameters inferred from MCMC for the same model gives
165 some indication of the importance of this uncertainty in the dynamic programming.

166 *The Gaussian Process model*

167 The core difference for our purpose between the estimated GP and the estimated parametric
168 models is that the estimated GP model is defined explicitly in reference to the observed data.
169 As a result, uncertainty arises in the GP model not only from uncertainty in the parameters,
170 but is also increases in regions farther from the observed states, such as low population sizes
171 in the example illustrated here. The estimated parametric models, by contrast, are completely
172 specified by the parameters.

173 The use of GPs to characterize dynamical systems is relatively new (Kocijan et al. 2005), and
174 was first introduced in the context ecological modeling and fisheries management in Munch et
175 al. (2005b). GP models have subsequently been used to test for the presence of Allee effects
176 (Sugeno and Munch 2013a), estimate the maximum reproductive rate (Sugeno and Munch
177 2013b), determine temporal variation in food availability (Sigourney et al. 2012), and provide a
178 basis for identifying model-misspecification (Thorson et al. 2014). An accessible and thorough
179 introduction to the formulation and use of GPs can be found in Rasmussen and Williams (2006).

180 A GP is a stochastic process for which any realization consisting of n points follows a multivariate
181 normal distribution of dimension n . To characterize the GP we need a mean function and a
182 covariance function. We proceed as follows.

183 As before, we assume that the data $X(t)$ are observed with process stochasticity around a mean
184 function $g(X(t))$

$$X(t+1) = g(X(t)) + \varepsilon, \quad (9)$$

where ε are IID normal random variables with zero-mean and variance σ^2 . Note that we have chosen to assume additive stochasticity. While we could consider log-normal stochasticity as in the parametric models, we make this choice to emphasize that the Gaussian process approach need not have structurally correct stochasticity to be effective.

In order to make predictions, we update the GP based on the observed set of transitions. To do so, we collect the time series of observed states into a vector of “current” states, $\mathbf{X}_{\text{obs}} = \{X(1), \dots, X(T-1)\}$ and a vector of “next” states $\mathbf{Y}_{\text{obs}} = \{X(2), \dots, X(T)\}$ where T is the time of the final observation. Conditional on these observations, the predicted next state, $g(X_p)$, for any given “current” state, X_p follows a normal distribution with mean E and variance C determined using the standard rules for conditioning in multivariate normals, i.e.

$$E = K(\mathbf{X}_p, \mathbf{X}_{\text{obs}}) (K(\mathbf{X}_{\text{obs}}, \mathbf{X}_{\text{obs}}) - \sigma \mathbb{I}_n)^{-1} \mathbf{Y}_{\text{obs}} \quad (10)$$

and

$$C = K(\mathbf{X}_p, \mathbf{X}_p) - K(\mathbf{X}_p, \mathbf{X}_{\text{obs}}) (K(\mathbf{X}_{\text{obs}}, \mathbf{X}_{\text{obs}}) - \sigma \mathbb{I})^{-1} K(\mathbf{X}_{\text{obs}}, \mathbf{X}_p) \quad (11)$$

Here \mathbb{I}_n is the n identity matrix (i.e. a matrix with ones down the diagonal and zeros elsewhere) and K is the ‘covariance kernel.’ The covariance kernel controls how much influence one observation has on another. In the present application we use the squared-exponential kernel which, when evaluated over a pair of vectors, say x and y , generates a covariance matrix whose i, j th element is given by

$$K_{i,j}(\mathbf{x}, \mathbf{y}) = \exp \left(\frac{-(x_i - y_j)^2}{2\ell^2} \right) \quad (12)$$

so that ℓ gives the characteristic length-scale over which correlation between two observations decays. See Rasmussen and Williams (2006) for other choices of covariance kernels and their

properties. Note that this simple formulation assumes a prior mean of zero. For the parameters we use inverse Gamma priors on both the length-scale ℓ and σ , thus for example

$$f(\ell; \alpha, \beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} \ell^{-\alpha-1} \exp\left(-\frac{\beta}{\ell}\right) \quad (13)$$

For the prior on ℓ , $\alpha = 10$ and $\beta = 10$. The prior on σ , $\alpha = 5$ and $\beta = 5$.

We use a Metropolis-Hastings Markov Chain Monte Carlo (Gelman et al. (2003)) to infer posterior distributions of the parameters of the GP (Figure S4, code in appendix). Since the posterior distributions differ substantially from the priors (Figure S4), most of the information in the posterior comes from the data rather than the prior belief.

The method of Gaussian Process Dynamic Programming (GPDP)

We derive the harvest policy from the estimated GP by inserting it into a SDP algorithm. Given the GP posteriors, we construct the transition matrix representing the probability of going to each state $X(t+1)$ given any current state $x(t)$ and any harvest $h(t)$ (See the function `gp_transition_matrix()` in the provided R package). Given this transition matrix, we use the same value iteration algorithm as in the parametric case to determine the optimal policy. In doing so, the uncertainty in the future state under the GP, $X(t+1)$, includes both process uncertainty (based on the estimation of σ) and structural uncertainty of the posterior collection of curves. Unlike the case of a parametric model, which is a single curve given the parameters, the GP remains a distribution of curves.

Results

Parametric and GP models for population dynamics

To ensure our results are robust to the choice of parameters, we will consider 96 different scenarios. To help better understand the process, we first describe in detail the results of a single scenario. All of the models fit the observed data rather closely and with relatively small uncertainty. In Figure 1, we show the posterior predictive curves. The training data of stock sizes observed




Figure 1: Points show the training data of stock-size over time. Curves show the expectations over the posterior step-ahead predictions based on each of the estimated models. (Observe that all models are fitting the data reasonably well.)

over time are points, overlaid with the step-ahead predictions of each estimated model using the parameters sampled from their posterior distributions. Compared to the true model most estimates appear to over-fit, predicting patterns that are actually due purely to stochasticity. Model selection criteria (Table 1) penalize more complex models and show a preference for the simpler Ricker model over the models with alternative stable states (Allen and Myers). Details on MCMC estimates for each model, traces, and posterior distributions can be found in the appendix.

	Allen	Ricker	Myers
DIC	50.75	50.45	50.41
AIC	-24.51	-30.13	-27.01
BIC	-17.75	-25.06	-20.25

Table 1: Model selection scores for several common criteria (DIC: Deviance Information Criterion, AIC: Akaike Information Criterion, BIC: Bayesian Information Criterion) all select the wrong model. As the true (Allen) model is not distinguishable from the simpler (Ricker) model in the region of the observed data, this error cannot be avoided regardless of the model choice criterion. This highlights the danger of model choice when the selected model will be used outside of the observed range of the data.

We show the mean inferred population dynamics of each model relative to the true model used to generate the data in Figure 2, predicting the relationship between observed population size (x-axis) to the population size after recruitment the following year. In addition to the raw data, the GP is conditioned on going through the point 0,0 without error. All parametric models also make this assumption. Conditioning on (0,0) is equivalent to making




Figure 2: The inferred Gaussian process compared to the true process and maximum-likelihood estimated (MLE) process. We show the expected value for the function f under each model. Plots show replicates drawn from the posterior distributions in order to convey uncertainty of the estimates. Note the MLE is a point estimate of parameters and so reflects no uncertainty in the distribution. The training data are also shown as black points. The GP is conditioned on $(0,0)$, shown as a pseudo-data point.

the assumption that the population is closed, so that once it hits 0 it stays at 0, despite the lack of any data in the observed sequence to justify this. This assumption illustrates how the GP can capture common-sense biology without having to assume more explicit details about the dynamics at low population numbers that have never been observed. If the population were not closed, one could repeat the entire analysis without this assumption. Unlike parametric models, the GP corresponds to a distribution of curves, of which this plot only shows the means. Uncertainty in the parameters of the GP (not shown) further widens the band of possible population sizes. In Figure S1 (see supplement), we show the performance of the models outside the observed training data.

Despite the similarities in model fits to the observed data, the policies inferred under each model differ widely (Figure 3). Policies are shown in terms of target escapement, $S(t) = x_t - h$. Under models such as this a constant escapement policy is expected to be optimal (Reed 1979), whereby population levels below a certain size S are unharvested, while above that size the harvest strategy aims to return the population to S . Whenever a model predicts that the population will not persist below a certain threshold, the optimal solution is to harvest the entire population immediately, resulting in an escapement $S = 0$, as seen in the true (correct form, exact parameters) model, the Allen model (correct form, estimated parameters) and the GP. Only the structurally correct model (Allen model) and the GP produce policies close to the true optimum policy.

In Figure 4, we show the consequences of managing 100 replicate realizations of the simulated fishery under policies derived from each model. The structurally correct model under-harvests, leaving the stock to vary around its unfished optimum. The structurally incorrect Ricker model over-harvests the population past the tipping point consistently, resulting in the immediate crash of the stock and thus leads to minimal long-term catch.

The results across replicate stochastic simulations are most easily compared by using the relative differences in net present value realized by each of the model (Figure 5). Although not perfect, the GPDP consistently realizes a value close to the optimal solution, and avoids ever driving the system across the tipping point, which results in the near-zero value cases in the parametric models.




Figure 3: The steady-state optimal policy (infinite boundary) calculated under each model. Policies are shown in terms of target escapement, $S(t)$, as under models such as this a constant escapement policy is expected to be optimal (Reed 1979). Several policies show a numerical jitter due to the discretization of states in the dynamic programming algorithm; doubling the number of grid points did not qualitatively change the results.

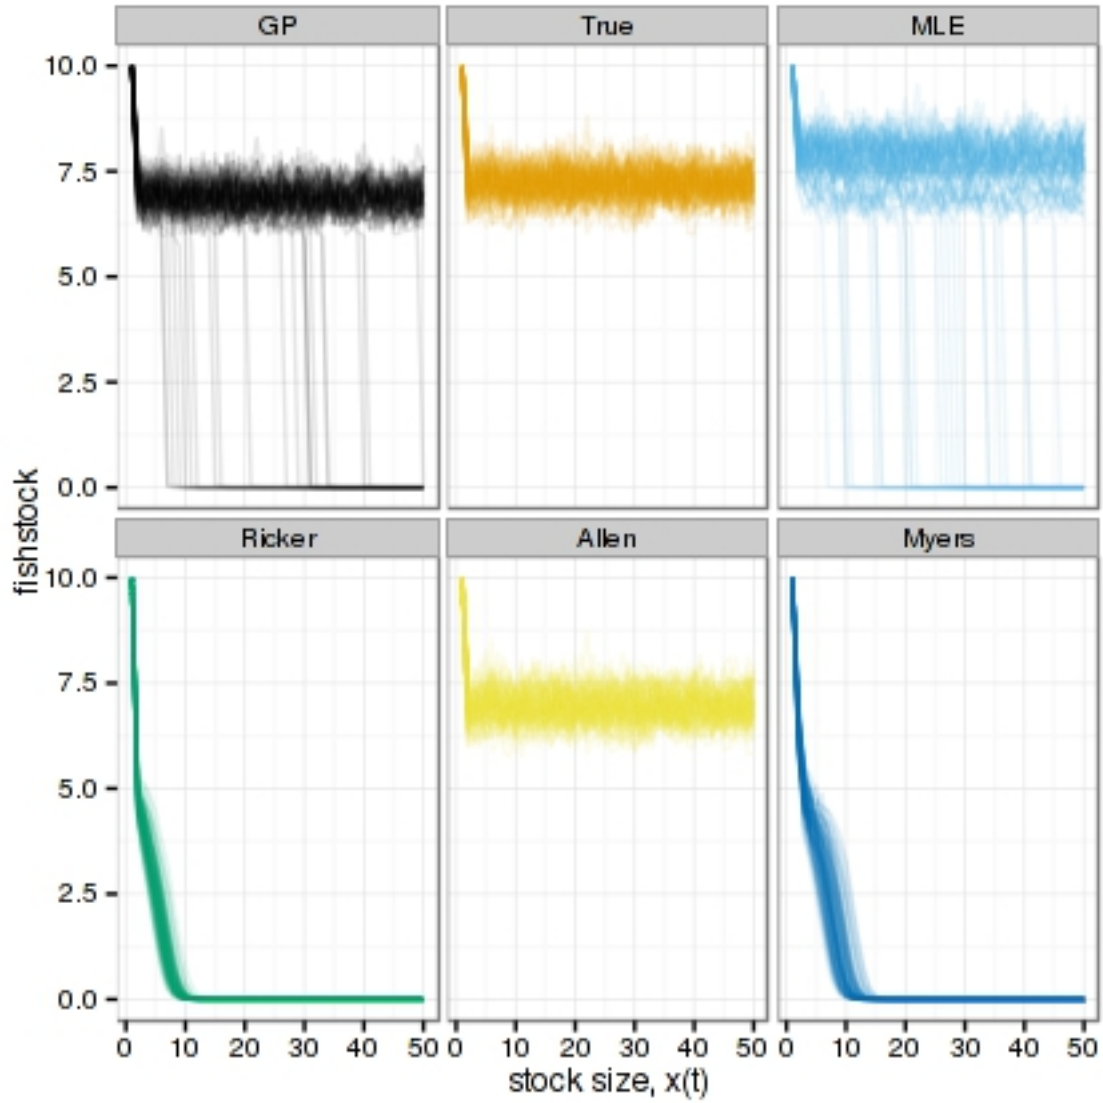


Figure 4: In the management context, GPDP outperforms approaches based on parametric models. We show 100 replicate simulations of the stock dynamics (Eqn 1) under the policies derived from each of the estimated models, as well as the policy based on the exact underlying model.

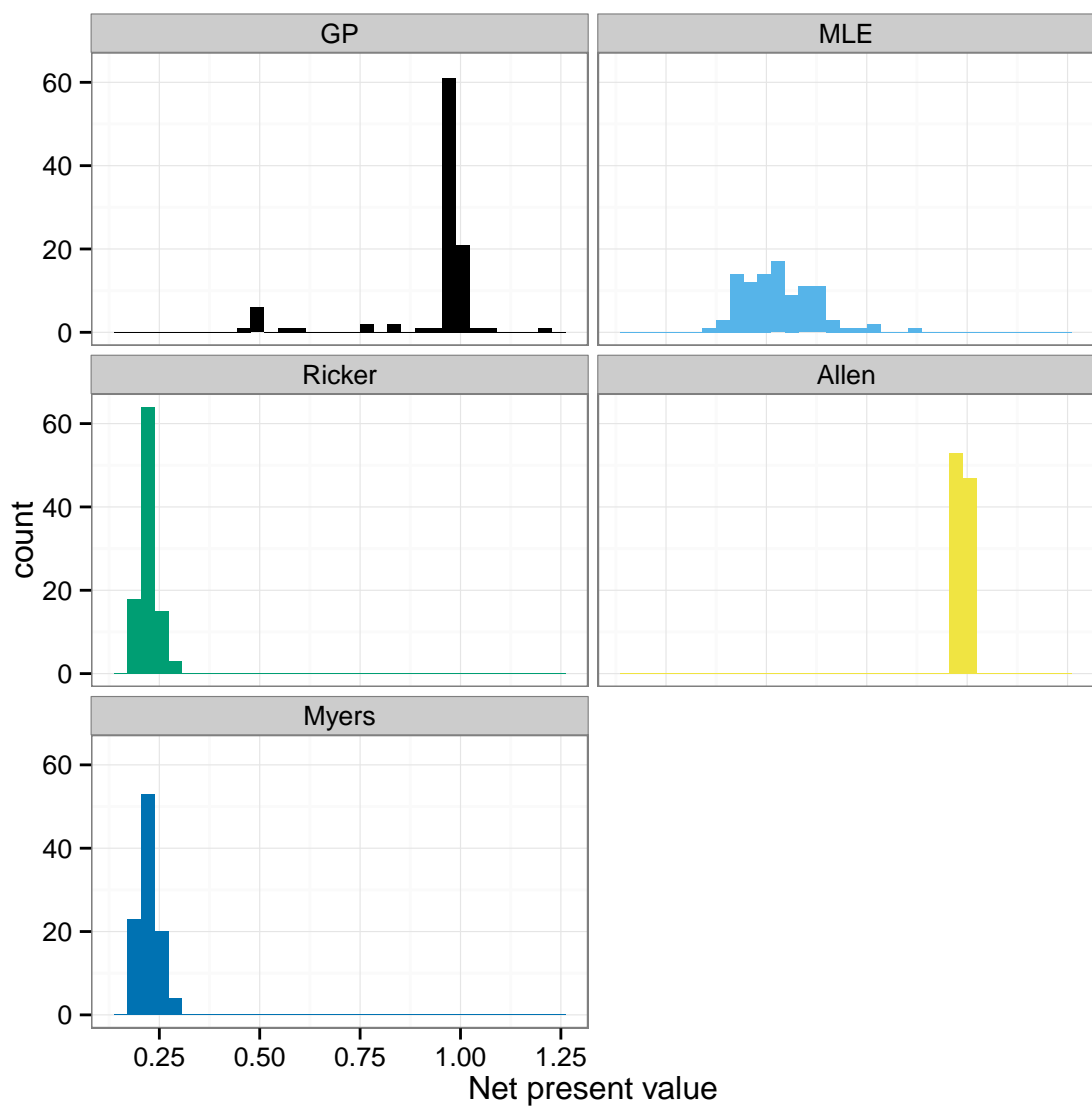


Figure 5: Histograms of the realized net present value of the fishery over a range of simulated data and resulting parameter estimates. For each data set, the three models are estimated as described above.

268 These results are not sensitive to the modeling details of the simulation. The GPDP estimate
 269 remains very close to the optimal solution (obtained by knowing the true model) across changes
 270 to the training simulation, scale of stochasticity, parameters or structure of the underlying
 271 model. In the Supplement, we consider both a Latin hypercube approach and a more focused
 272 investigation of the effects of the relative distance to the Allee threshold and the variance of
 273 process stochasticity.

274 Changing the intensity of the stochasticity or the distance between stable and unstable steady
 275 states does not impact the performance of the GP relative to the optimal solution obtained from
 276 the true model and true parameters. The parametric models are more sensitive to this difference.
 277 Large values of σ relative to the distance between the stable and unstable point increases the
 278 chance of a stochastic transition below the tipping point. More precisely, if we let L be the
 279 distance between the stable and unstable steady states, then the probability that fluctuations
 280 drive the population across the unstable steady state scales as $\exp\left(\frac{L^2}{\sigma^2}\right)$
 281 (see Gardiner (2009) or Mangel (2006) for the derivation). Thus, the impact of using a model
 282 that underestimates the risk of harvesting beyond the critical point is considerable, since
 283 this such a situation occurs more often. Conversely, with large enough distance between the
 284 optimal escapement and unstable steady state relative to σ , the chance of a transition becomes
 285 vanishingly small and all models can be estimated near-optimally. Models that underestimate
 286 the cost incurred by population sizes fluctuating significantly below the optimal escapement level
 287 will not perform poorly as long as those fluctuations are sufficiently small.

288 The GPDP is only weakly influenced by increasing stochasticity or increasing Allee effects over
 289 much of the range (Figure 6). Larger σ or higher Allee levels make even the optimal solution
 290 without any model or parameter uncertainty unable to harvest the population effectively (e.g. the
 291 stochasticity is large enough to violate the self-sustaining criterion of Reed (1979)).

292 **Discussion**

293 Simple, mechanistically motivated models offer the potential to increase our basic understanding
 294 of ecological processes (Geritz and Kisdi 2012, Cuddington et al. 2013). But such models can be




Figure 6: The effect of increasing noise or decreasing Allee threshold levels on the net present value of the fishery when managed under the GPDP, relative to managing under the true model (with known parameters). Other than the focal parameter (stochasticity, Allee threshold), other parameters are held fixed as above to illustrate this effect.

both inaccurate and misleading when used in a decision making framework. In this paper we tackled two aspects of uncertainty that are common to many ecological decision-making problems and fundamentally challenging to existing approaches that largely rely on parametric models:

1. We do not know the correct models for ecological systems.
2. We have limited data from which to estimate the model.

We have illustrated how the use of non-parametric methods provides more reliable solutions in the sequential decision-making problem.

Traditional model-choice approaches can be positively misleading.

Our results illustrate that model-choice approaches can be absolutely misleading – by providing support to models that cannot capture tipping point dynamics because they have fewer parameters and the data are far from the tipping point. That is, when the data come from around the stable steady state, all the parametric models are approximately linear and approximately identical. Thus, it is intuitive that all model selection methods choose the simplest model. In a complex

308 world, the result is suboptimal. But in a world that might contain tipping points, the result
309 could be disastrous.

310 Many managers both in fisheries and beyond face a similar situation: they have not observed the
311 population dynamics at all possible densities. A lack of comprehensive data at all population
312 sizes, combined with the inability to formulate accurate population models for low population
313 sizes in the absence of data, makes this situation the rule more than the exception. Relying on
314 parametric models and model choice processes that favor simplicity ignores this basic reality. For
315 a long time, Carl Walters (e.g. Walters and Hilborn 1978) has argued that if we began by fishing
316 any newly exploited population down to very low levels and then let it recover, we would be
317 much better at estimating population dynamics and thus predicting the optimal harvest levels.
318 While certainly true, this presents a rather risky policy in the face of potential tipping points.
319 The GPDP offers a risk-adverse alternative.

320 *GPDP population dynamics capture larger uncertainty in regions where the data are poor*

321 Parametric models perform most poorly when we seek a management strategy outside the range
322 of the observed data. The GPDP, in contrast, leads to a predictive model that expresses a great
323 deal of uncertainty about the probable dynamics *outside* the range of the observed data, while
324 retaining very good predictive accuracy *inside* the range. The management policy based on by
325 the GPDP balances uncertainty outside the range of the observed data against the immediate
326 value of the harvest, and acts to stabilize the population dynamics in a region of state space in
327 which the predictions are reliably reflected by the data.

328 Such problems are ubiquitous across ecological decision-making and conservation where the
329 greatest concerns involve decisions that lead to population sizes that have never been observed
330 and for which we do not know the response – whether this is the collapse of a fishery, the spread
331 of an invasive, or the loss of habitat.

332 *The role of the prior*

333 Outside of the observed range of the data, the GP reverts to the prior, and consequently the
334 choice of the prior can also play a significant role in determining the optimal policy. In the

examples shown here we have selected a prior that is both relatively uninformative (due to the broad priors placed on its parameters ℓ and σ and simple (the choice of our covariance function, Eqns 12 and 13). In practice, these should be chosen to confer particular biological properties. In principle, this may allow a manager to improve the performance of the GPDP by adding detail as is justified. For instance, it would be possible to use a linear or a Ricker-shaped mean in the prior without making the much stronger assumption that the Ricker is the structurally correct model (Sugeno and Munch 2013a). One fruitful avenue of future research is identifying criteria to ensure the prior and the reward function are chosen appropriately for the problem at hand.

Acknowledgments

This work was partially supported by NOAA-IAM grant to SM and Alec McCall and administered through the Center for Stock Assessment Research, a partnership between the University of California Santa Cruz and the Fisheries Ecology Division, Southwest Fisheries Science Center, Santa Cruz, CA and by NSF grant EF-0924195 to MM and NSF grant DBI-1306697 to CB.

Allen, D., and K. Tanner. 2005. Infusing active learning into the large-enrollment biology class: seven strategies, from the simple to complex. *Cell biology education* 4:262–8.

Athanassoglou, S., and A. Xepapadeas. 2012. Pollution control with uncertain stock dynamics: When, and how, to be precautionous. *Journal of Environmental Economics and Management* 63:304–320.

Brozović, N., and W. Schlenker. 2011. Optimal management of an ecosystem with an unknown threshold. *Ecological Economics*:1–14.

Burnham, K. P., and D. R. Anderson. 2002. *Model Selection and Multi-Model Inference*. Page 496. Springer.

Clark, C. W. 1976. *Mathematical Bioeconomics*. WileyNew York.

Clark, C. W. 2009. *Mathematical Bioeconomics*. WileyNew York.

Clark, C. W., and G. P. Kirkwood. 1986. On uncertain renewable resource stocks: Optimal harvest policies and the value of stock surveys. *Journal of Environmental Economics and Management* 13:235–244.

362 Clark, C. W., and M. Mangel. 2000. Dynamic state variable models in ecology. Oxford University
363 PressOxford.

364 Courchamp, F., L. Berec, and J. Gascoigne. 2008. Allee Effects in Ecology and Conservation.
365 Page 256. Oxford University Press, USA.

366 Cressie, N., C. a Calder, J. S. Clark, J. M. Ver Hoef, and C. K. Wikle. 2009. Accounting
367 for uncertainty in ecological analysis: the strengths and limitations of hierarchical statistical
368 modeling. *Ecological Applications* 19:553–70.

369 Cuddington, K. M., M. Fortin, and L. Gerber. 2013. Process-based models are required to
370 manage ecological systems in a changing world. *Ecosphere* 4:1–12.

371 Fischer, J., G. D. Peterson, T. a Gardner, L. J. Gordon, I. Fazey, T. Elmqvist, A. Felton, C.
372 Folke, and S. Dovers. 2009. Integrating resilience thinking and optimisation for conservation.
373 *Trends in ecology & evolution* 24:549–54.

374 Gardiner, C. 2009. Stochastic Methods: A Handbook for the Natural and Social Sciences
375 (Springer Series in Synergetics). Page 447. Springer.

376 Gelman, A., J. B. Carlin, H. S. Stern, and D. B. Rubin. 2003. Bayesian Data Analysis. 2nd
377 editions. Chapman; Hall/CRC.

378 Geritz, S. A. H., and E. Kisdi. 2012. Mathematical ecology: why mechanistic models? *Journal*
379 *of mathematical biology* 65:1411–5.

380 Gordon, H. 1954. The economic theory of a common-property resource: the fishery. *The Journal*
381 *of Political Economy* 62:124–142.

382 Hilborn, R., and M. Mangel. 1997. The Ecological Detective: Confronting Models with data.
383 Page 330. Princeton University Press.

384 Hughes, T. P., C. Linares, V. Dakos, I. a van de Leemput, and E. H. van Nes. 2013. Living
385 dangerously on borrowed time during slow, unrecognized regime shifts. *Trends in ecology &*
386 *evolution* 28:149–55.

387 Kocijan, J., A. Girard, B. Banko, and R. Murray-Smith. 2005. Dynamic systems identification
388 with Gaussian processes. *Mathematical and Computer Modelling of Dynamical Systems* 11:411–

389 424.

390 Ludwig, D., and C. J. Walters. 1982. Optimal harvesting with imprecise parameter estimates.
391 Ecological Modelling 14:273–292.

392 Mangel, M. 2006. The Theoretical Biologist’s Toolbox: Quantitative Methods for Ecology and
393 Evolutionary Biology. Cambridge University Press.

394 Mangel, M. 2014. Stochastic Dynamic Programming Illuminates the Link Between Environment.
395 Bulletin of Mathematical Biology in press.

396 Mangel, M., and C. W. Clark. 1988. Dynamic Modeling in Behavioral Ecology. (J. Krebs and T.
397 Clutton-Brock, Eds.). Princeton University PressPrinceton.

398 Mangel, M., O. Fiksen, and J. Giske. 2001. Theoretical and statistical models in natural
399 resource management and research. Pages 57–71 *in* T. M. Shenk and A. B. Franklin, editors.
400 Modeling in natural resource management, development, interpretation and application. Island
401 PressWashington DC.

402 Marescot, L., G. Chapron, I. Chadès, P. L. Fackler, C. Duchamp, E. Marboutin, and O. Gimenez.
403 2013. Complex decisions made simple: a primer on stochastic dynamic programming. Methods
404 in Ecology and Evolution:n/a–n/a.

405 May, R. M., J. R. Beddington, C. W. Clark, S. J. Holt, and R. M. Laws. 1979. Management of
406 multispecies fisheries. Science (New York, N.Y.) 205:267–77.

407 McAllister, M. 1998. Bayesian stock assessment: a review and example application using the
408 logistic model. ICES Journal of Marine Science 55:1031–1060.

409 Munch, S. B., A. Kottas, and M. Mangel. 2005a. Bayesian nonparametric analysis of stock-
410 recruitment relationships. Canadian Journal of Fisheries and Aquatic Sciences 62:1808–1821.

411 Munch, S. B., M. L. Snover, G. M. Watters, and M. Mangel. 2005b. A unified treatment of
412 top-down and bottom-up control of reproduction in populations. Ecology Letters 8:691–695.

413 Myers, R. A., N. J. Barrowman, J. A. Hutchings, and A. a Rosenberg. 1995. Population dynamics
414 of exploited fish stocks at low population levels. Science (New York, N.Y.) 269:1106–8.

415 Polasky, S., S. R. Carpenter, C. Folke, and B. Keeler. 2011. Decision-making under great

416 uncertainty: environmental management in an era of global change. *Trends in ecology &*
417 *evolution*:1–7.

418 R Core Team. 2013. R: A Language and Environment for Statistical Computing. R Foundation
419 for Statistical Computing Vienna, Austria.

420 Rasmussen, C. E., and C. K. I. Williams. 2006. *Gaussian Processes for Machine Learning*.
421 (Thomas Dietterich, Ed.). MIT Press, Boston.

422 Reed, W. J. 1979. Optimal escapement levels in stochastic and deterministic harvesting models.
423 *Journal of Environmental Economics and Management* 6:350–363.

424 Roughgarden, J. E., and F. Smith. 1996. Why fisheries collapse and what to do about it.
425 *Proceedings of the National Academy of Sciences of the United States of America* 93:5078.

426 Schapaugh, A. W., and A. J. Tyre. 2013. Accounting for parametric uncertainty in Markov
427 decision processes. *Ecological Modelling* 254:15–21.

428 Scheffer, M., J. Bascompte, W. A. Brock, V. Brovkin, S. R. Carpenter, V. Dakos, H. Held, E.
429 H. van Nes, M. Rietkerk, and G. Sugihara. 2009. Early-warning signals for critical transitions.
430 *Nature* 461:53–9.

431 Scheffer, M., S. R. Carpenter, J. A. Foley, C. Folke, and B. Walker. 2001. Catastrophic shifts in
432 ecosystems. *Nature* 413:591–6.

433 Sethi, G., C. Costello, A. Fisher, M. Hanemann, and L. Karp. 2005. Fishery management under
434 multiple uncertainty. *Journal of Environmental Economics and Management* 50:300–318.

435 Sigourney, D. B., S. B. Munch, and B. H. Letcher. 2012. Combining a Bayesian nonparametric
436 method with a hierarchical framework to estimate individual and temporal variation in growth.
437 *Ecological Modelling* 247:125–134.

438 Su, Y.-S., and Masanao Yajima. 2013. R2jags: A Package for Running jags from R.

439 Sugeno, M., and S. B. Munch. 2013a. A semiparametric Bayesian method for detecting Allee
440 effects. *Ecology* 94:1196–1204.

441 Sugeno, M., and S. B. Munch. 2013b. A semiparametric Bayesian approach to estimating
442 maximum reproductive rates at low population sizes. *Ecological applications* : a publication of

- 443 the Ecological Society of America 23:699–709.
- 444 Thorson, J. T., K. Ono, and S. B. Munch. 2014. A Bayesian approach to identifying and
445 compensating for model misspecification in population models. *Ecology* 95:329–41.
- 446 Walters, C. J., and R. Hilborn. 1978. Ecological Optimization and Adaptive Management.
447 *Annual Review of Ecology and Systematics* 9:157–188.
- 448 Weitzman, M. L. 2013. A Precautionary Tale of Uncertain Tail Fattening. *Environmental and*
449 *Resource Economics* 55:159–173.
- 450 Williams, B. K. 2001. Uncertainty , learning , and the optimal management of wildlife. *Environ-*
451 *mental and Ecological Statistics* 8:269–288.